


Homework 7

Theory Question

Within the neural network, first the image is processed into 4 different scales: the original image, $1/2$ resolution, $1/4$ resolution and $1/8$ resolution. After passing through a VGG-19 like structure (but only the first 16 layers, with modifications to have the resulting sizes not be too small), there are 4 convolutional blocks of varying depths. Each of the image resolutions is passed through these 4 blocks, for the outputs that have a size smaller than 32 by 32 , we upsample them to reach size 32 by 32 . Then all of these outputs are concatenated together, we pass this result through one last convolutional layer that downsizes the number of channels to 512 but keeps the image resolution the same (32 by 32). We then calculate the channel normalization parameters for these channels, which we then use in a similar way to AdaIN.

This method might work best on images that contain most of the features within the first few scales ($1/2$, $1/4$ and $1/8$), as we are explicitly running VGG on these. So there could be some information loss for the textures that are detected at higher scales. Being able to put the images through one more block will allow us to get a little bit more features, and then the concatenation and subsequent convolution allows us to relate information from all scales as well as their features to enhance our classifier and also make sure that there is no redundant information from them as dictated by the weights of this final layer. This would not perform well in images where there are small objects, or where the textures all look similar at the scales we work on.

Task




(a) Example for "cloudy" class.



(b) Example for "rain" class.



(c) Example for "shine" class.



(d) Example for "sunrise" class.

Figure 1: Example images for each class.

LBP

First we convert the image to the HSI representation. We do this by following the algorithm to convert from RGB to HSI from Prof. Kak's book. Additionally, we round up and divide the H channel by 2, this is done to be able to use `cv2.resize` on the resulting matrix. After this, we loop through the image and use `BitVector` to encode the pattern. We use $R = 2$ and $P = 12$ for our LBP extraction. We found that this parameters worked a bit better compared to $R = 1$ and $P = 8$. After this, we convert the output to an array and divide by the sum to get values from 0 to 1. These features are what we use to train the SVM on the training dataset and then get the confusion matrix and the accuracy on the testing dataset.

Gram-matrix

For all models we follow the same procedure. First each image is resized to $(256, 256)$ and we get the output for the models and then reshape it, then multiply it by its transpose, which is symmetric so we take the upper triangular part. Since the final matrix is of size

(512, 512) or in the case of Resnet coarse (1024, 1024), we flatten the array, and take the first 2048 values. The reason for 2048 is that since we only took the upper triangular matrix, the dimensionality is half of 2048, 1024. This is to compare how well this does against the AdaIN features in the following section. This feature vector of size 2048 is what is used to train the SVM.

AdaIN features

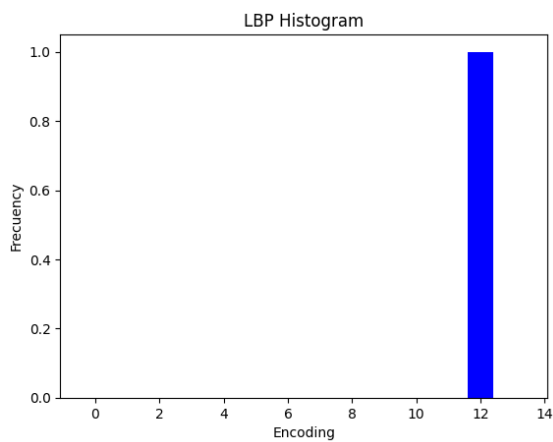
We use VGG model for this part. We take the output of shape (512, 32, 32) and calculate the mean and standard deviation per channel resulting in 2 arrays of size 512. We concatenate both and use that array of length 1024 to train our SVM.

Results


LBP

There was one issue we found when looking at the images. Some of the classes, mainly "cloudy" contained grayscale images, the problem with this is that when we get the HSI representation, the hue channel is all 1. This is due to the calculation of the maximum and minimum RGB value, but in grayscale they are all the same. This can be seen in Fig. 2a. This could potentially cause issues when classifying using the SVM.


After training the SVM, it performed poorly. Shown in Fig. 3, we can see that only the "sunrise" class performed well, and that is because it mostly classified everything as "sunrise". This could be explained by the fact that "sunrise" class contains the most amount of images. The class it struggled with the most is "shine", and it mostly classified all of them as "sunrise". There is quite a bit of overlap between "shine" and "sunrise" since they both feature the sun, there are also clouds which explains why part of the "cloudy" images were classified as "sunrise" and there is also some overlap between "rain" and "cloudy" as they both feature clouds. Overall, the accuracy for this classifier on the LBP descriptors was 40% on the testing dataset.



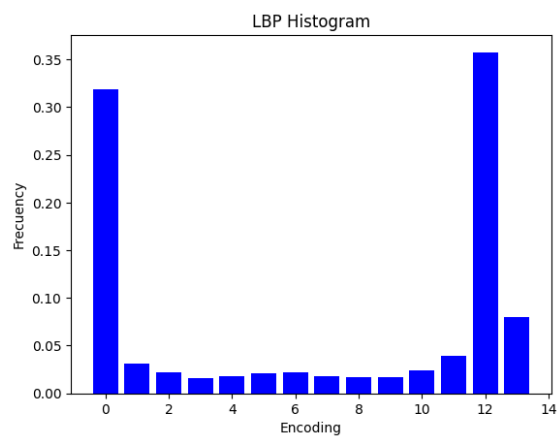
(a) LBP descriptor vector for "cloudy" class.



(b) LBP descriptor vector for "rain" class.



(c) LBP descriptor vector for "shine" class.



(d) LBP descriptor vector for "sunrise" class.

Figure 2: LBP descriptor vector for each class.

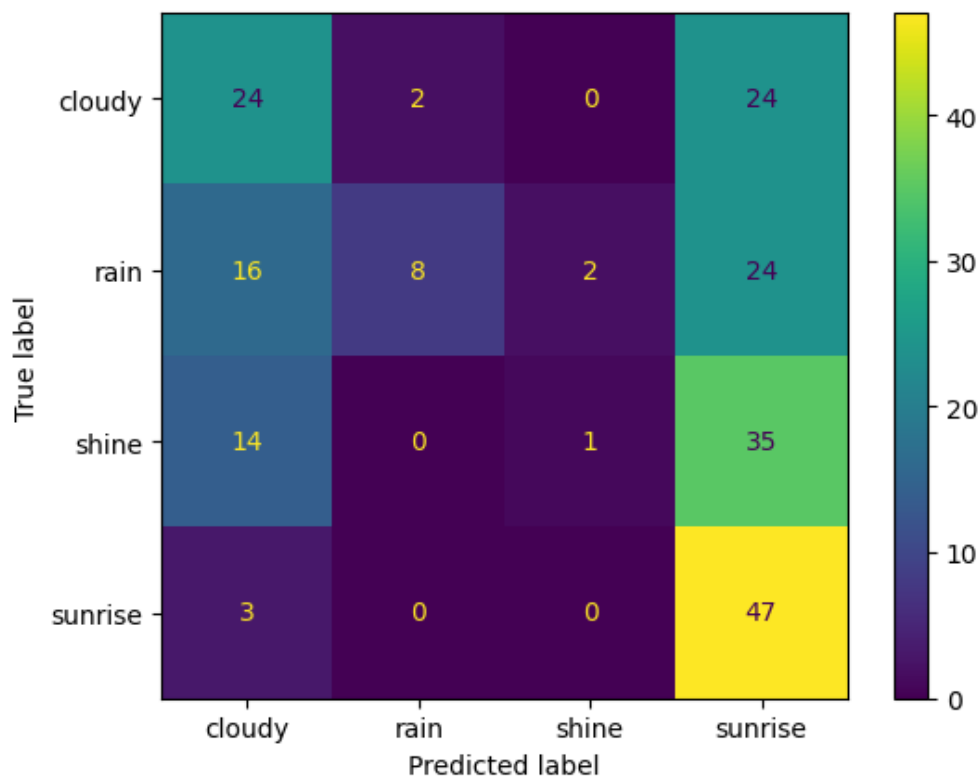


Figure 3: Confusion matrix for the LBP descriptors



(a) Image correctly classified as "rain" class.




(b) Image incorrectly classified as "cloudy" class, the true label is "shine".

Figure 4: Correct and incorrect classification for LBP.


Gram-Matrix

Resnet Coarse For Resnet coarse, it is hard to spot many differences in the Gram matrix representation, as shown in Fig. 5. We no longer have the issue with grayscale images here, as we use the RGB representation instead of HSI.


After training the SVM, with a confusion matrix shown in Fig. 6, we notice an improvement in the performance compared to LBP. In this case, we correctly classified all "sunrise" images, and most of the "cloudy" images got classified correctly. The classes that this classifier struggled on the most was "rain" and "shine", with both of them being mostly misclassified as "cloudy", this could be due to the fact that both of those classes have images that have clouds in them. Overall, this classifier achieved an accuracy of 79.5% on the testing set.




(a) Gram matrix from Resnet coarse for "cloudy" class.



(b) Gram matrix from Resnet coarse for "rain" class.



(c) Gram matrix from Resnet coarse for "shine" class.



(d) Gram matrix from Resnet coarse for "sunrise" class.

Figure 5: Gram matrix from Resnet coarse for each class.




Figure 6: Confusion matrix for the Resnet coarse Gram matrix



(a) Image correctly classified as "cloudy" class.




(b) Image incorrectly classified as "sunrise" class, the true label is "rain".


Figure 7: Correct and incorrect classification for Resnet coarse.

Resnet Fine Now we can see some differences in the Gram matrix representations, as shown in Fig. 8. After training the SVM, notice that the confusion matrix, Fig. 9, shows significant improvements over Resnet coarse. Most of the images have been correctly classified, the only class that shows some struggles is "shine", with some classified as "cloudy"


and some classified as "sunrise", this could be explained because, in a similar manner to Resnet coarse, all of them contain images with clouds. The accuracy on the testing set for this classifier is 91.5%.



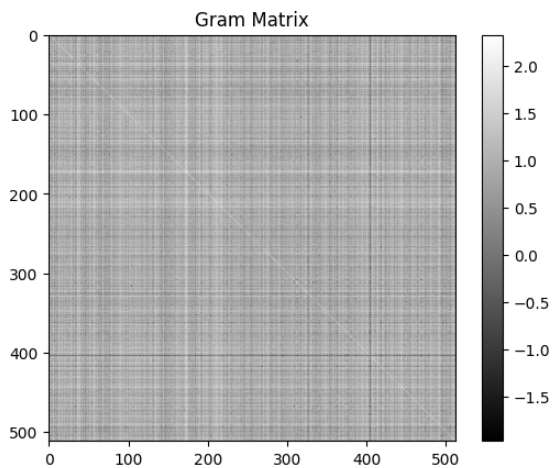
(a) Gram matrix from Resnet fine for "cloudy" class.



(b) Gram matrix from Resnet fine for "rain" class.



(c) Gram matrix from Resnet fine for "shine" class.



(d) Gram matrix from Resnet fine for "sunrise" class.

Figure 8: Gram matrix from Resnet fine for each class.




Figure 9: Confusion matrix for the Resnet fine Gram matrix



(a) Image correctly classified as "sunrise" class.



(b) Image incorrectly classified as "sunrise" class, the true label is "rain".

Figure 10: Correct and incorrect classification for Resnet fine.

VGG The differences in the Gram matrix representations, as shown in Fig. 11 are very pronounced now. The SVM classifier does well on both "sunrise" and "rain" but is not as accurate on the other classes, as can be seen in Fig. 12. The performance of this model is worse than Resnet fine but better than Resnet coarse. There are many images misclassified

as "sunrise", possibly due to the fact that they have the sun. The overall accuracy for this model on the testing set is 89%.

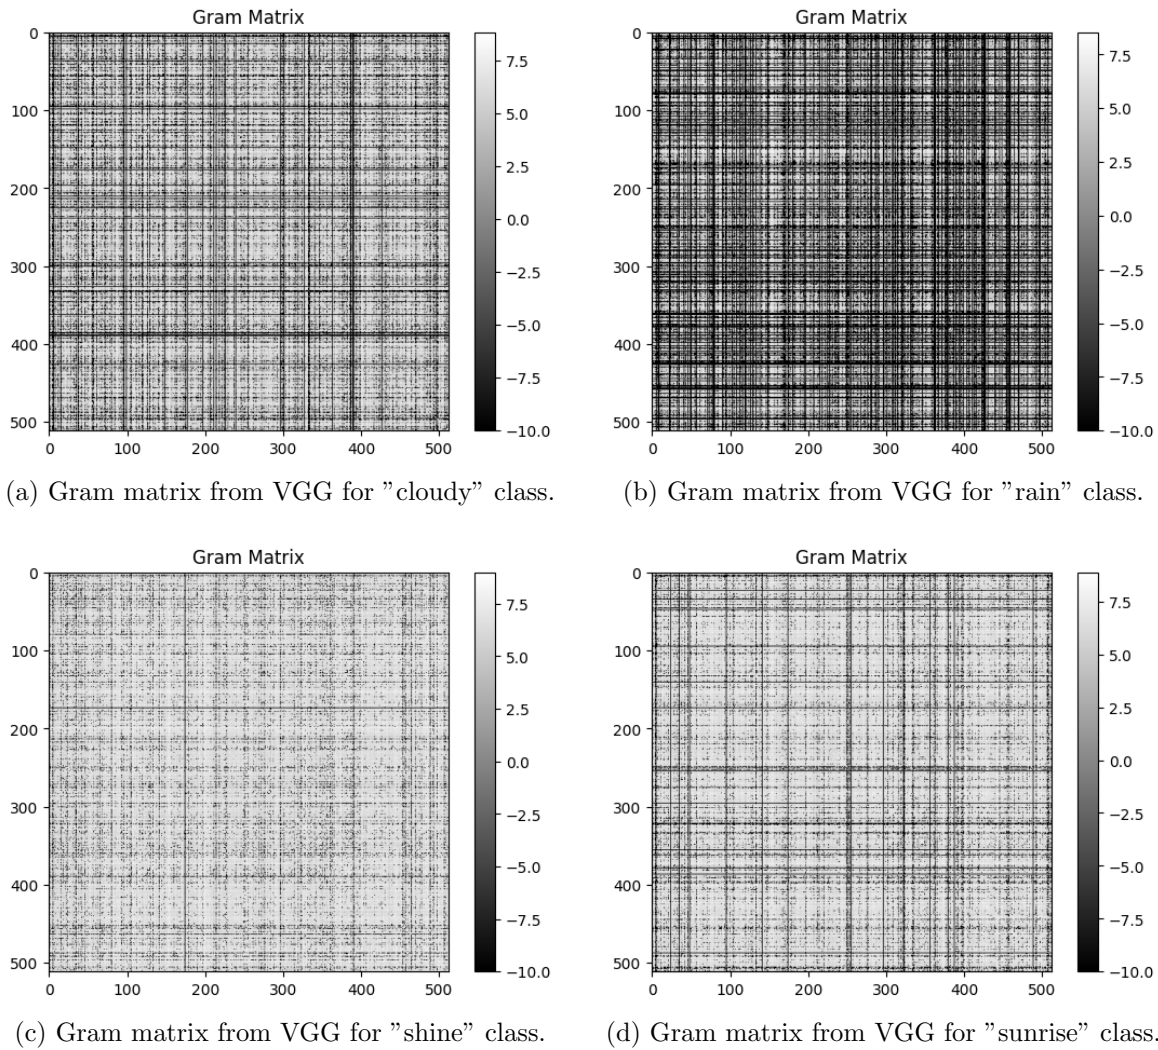


Figure 11: Gram matrix from Resnet fine for each class.





Figure 12: Confusion matrix for the VGG Gram matrix



(a) Image correctly classified as "shine" class.



(b) Image incorrectly classified as "sunrise" class, the true label is "cloudy".

Figure 13: Correct and incorrect classification for Resnet coarse.

Extra credit: AdaIN

AdaIN features performed the best out of all classifiers. As can be seen in Fig. 14, it performed extremely well on "cloudy", "rain" and "sunrise". Only the "shine" class had 4

images misclassified as both "sunrise" and "cloudy" due to the sun and the clouds in the images respectively. The total accuracy on the testing set is 97.5%.




Figure 14: Confusion matrix for the AdaIN features

Source code

```

1  import cv2
2  import numpy as np
3  import matplotlib.pyplot as plt
4  import BitVector
5  import os
6  from vgg_and_resnet import *
7  from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
8  from sklearn import svm
9
10 def get_hsl(image):
11     image = image/255.
12     # image is bgr, opencv default
13     # start by splitting image into blue green red
14     blue, green, red = image[...,0], image[...,1], image[...,2]
15     # calculate the max and min per position
16     Cmax_arg = np.argmax(image,axis=2)
17     Cmax = np.max(image,axis=2)
18     Cmin = np.min(image,axis=2)
19     # get the delta

```

```

20     delta = Cmax - Cmin
21     # now start populating the hsv matrices
22     v = Cmax
23     h = np.zeros_like(v)
24     s = np.zeros_like(v)
25     s[v != 0] = delta[v != 0] / v[v != 0]
26
27     # was not entirely sure how to vectorize this so I did it with a loop
28     for j in range(h.shape[0]):
29         for i in range(h.shape[1]):
30             if delta[j,i] != 0:
31                 # for blue
32                 if Cmax_arg[j,i] == 0:
33                     h[j,i] = ((60 * (red[j,i] - green[j,i]) / delta[j,i])
+ 240) % 360
34
35                 # for green
36                 elif Cmax_arg[j,i] == 1:
37                     h[j,i] = ((60 * (blue[j,i] - red[j,i]) / delta[j,i]) +
120) % 360
38
39                 # and red
40                 else:
41                     h[j,i] = ((60 * (green[j,i] - blue[j,i]) / delta[j,i])
+ 360) % 360
42     # the hue channel is halved and made to be integer so I can use the
opencv resize function
43     # this is the same process opencv does to get the hsi value
44     return np.round(h/2).astype(np.uint8), s, i
45
46 def get_lbp_descriptor(image, P=8, R=1):
47     # we use bitvector and professor Kak's implementation from his book
48     # get the hsl representation then resize
49     img_h, _, _ = get_hsl(image)
50     img = cv2.resize(img_h, (64,64))
51     eps = 1e-5
52     # this is based on professor Kak's implementation
53     # since its a square image, height and width are the same
54     r_max = 64 - R
55     # we start with a dictionary but then convert it to a numpy array with
relative values
56     lbp_hist = {t:0 for t in range(P+2)}
57     # we set the arrays that contain the sin and cos for the neighbors
58     pp = np.arange(P)
59     pp_x = R*np.cos(2*np.pi*pp/P)
60     pp_y = R*np.sin(2*np.pi*pp/P)
61     pp_x[pp_x < eps] = 0
62     pp_y[pp_y < eps] = 0
63
64     for j in range(R,r_max):
65         for i in range(R,r_max):
66
67             pattern = []
68             x_ = i + pp_x

```

```

69     y_ = j + pp_y
70     for p_x, p_y in zip(x_, y_):
71         x_base, y_base = int(p_x), int(p_y)
72         x_delta, y_delta = p_x - x_base, p_y - y_base
73         if (x_delta < eps) and (y_delta < eps):
74             image_p = float(img[x_base][y_base])
75         elif (y_delta < eps):
76             image_p = (1 - x_delta) * img[x_base][y_base] +
x_delta * img[x_base+1][y_base]
77         elif (x_delta < eps):
78             image_p = (1 - y_delta) * img[x_base][y_base] +
y_delta * img[x_base][y_base+1]
79         else:
80             image_p = (1 - x_delta)*(1 - y_delta)*img[x_base][
y_base] + (1-x_delta)*y_delta*img[x_base][y_base + 1] + x_delta*y_delta
*img[x_base+1][y_base+1] + x_delta*(1-y_delta)*img[x_base+1][y_base]
81         if image_p >= img[j][i]:
82             pattern.append(1)
83         else:
84             pattern.append(0)
85     bitv = BitVector.BitVector( bitlist = pattern )
86     intvals_for_circular_shifts = [int(bitv << 1) for _ in range(P
)]
87     minbitval = BitVector.BitVector( intVal = min(
intvals_for_circular_shifts), size = P )
88     bvrns = minbitval.runs()
89     if len(bvrns) > 2:
90         lbp_hist[P+1] += 1
91     elif len(bvrns) == 1 and bvrns[0][0] == "1":
92         lbp_hist[P] += 1
93     elif len(bvrns) == 1 and bvrns[0][0] == "0":
94         lbp_hist[0] += 1
95     else:
96         lbp_hist[len(bvrns[1])] += 1
97     # now we get the numpy array from the values of the dictionary and
normalize the array
98     lbp_hist = np.array(list(lbp_hist.values()))
99
100     return lbp_hist/lbp_hist.sum()
101
102 def get_gram_descriptor(image, model, coarse=None):
103     # we use this function to calculate the gram matrix for the image
104     # we start by resizing
105     image = cv2.resize(image, (256, 256))
106     # since we have resnet coarse and fine and vgg we use the variable
coarse to handle each case
107     # coarse = None is for vgg
108     # coarse = True is for coarse resnet
109     # coarse = False is for fine resnet
110     if coarse == None:
111         img_features = model(image)
112         # flatten the array
113         img_features = img_features.reshape((img_features.shape[0],
img_features.shape[1]*img_features.shape[2]))

```

```

114         # get the gram matrix
115         gmatrix = img_features @ img_features.T
116         # normalize it
117         g_max = gmatrix.max()
118         # we are only interested in the upper triangular
119         return np.triu(gmatrix)/g_max, gmatrix
120     # we repeat this process for the other 2 models
121     elif coarse:
122         img_features, _ = model(image)
123         img_features = img_features.reshape((img_features.shape[0],
124 img_features.shape[1]*img_features.shape[2]))
125         gmatrix = img_features @ img_features.T
126         g_max = gmatrix.max()
127         return np.triu(gmatrix)/g_max, gmatrix
128     elif not coarse:
129         _, img_features = model(image)
130         img_features = img_features.reshape((img_features.shape[0],
131 img_features.shape[1]*img_features.shape[2]))
132         gmatrix = img_features @ img_features.T
133         g_max = gmatrix.max()
134         return np.triu(gmatrix)/g_max, gmatrix
135
136 def get_adain_features(image, model):
137     # we start by resizing the image to 256 256
138     image = cv2.resize(image,(256,256))
139     # get the output
140     img_features = model(image)
141     # flatten the output
142     img_features = img_features.reshape((img_features.shape[0],
143 img_features.shape[1]*img_features.shape[2]))
144     # get the mean and standard deviation per channel (axis = 1)
145     img_features_mean = img_features.mean(axis=1)
146     img_features_std = img_features.std(axis=1)
147     # concatenate both mean and standard deviation,
148     # this is our feature vector now
149     adain_features = np.hstack((img_features_mean,img_features_std))
150     return adain_features
151
152 def get_class_lbp(path,P=8, R=1, class_name="cloudy"):
153     # I use this function to loop through the class images and get a list
154     # of all the descriptor vectors corresponding to that class
155     file_list = [x for x in os.listdir(path) if (class_name in x and x.
156 endswith(".jpg"))]
157     lbp_feat = []
158     test_names = []
159     for idx in range(len(file_list)):
160         # there is a try here since some of the images cannot be opened as
161         # they appear to be gifs
162         try:
163             img = cv2.imread(os.path.join(path,file_list[idx]))
164             lbp_desc = get_lbp_descriptor(img, P=P, R=R)
165             lbp_feat.append(lbp_desc)
166             # I also append the names here to know which are the
167             # misclassified images

```

```

161         test_names.append(file_list[idx])
162     except:
163
164         print(file_list[idx])
165
166     return np.array(lbp_feat), test_names
167
168 def get_class_gram(path,model, class_name="cloudy", coarse=None):
169     # I use this function to loop through the class images and get a list
170     # of all the gram matrices that correspond to one class
171     # similar to the lbp function
172     file_list = [x for x in os.listdir(path) if (class_name in x and x.
173     endswith(".jpg"))]
174     gram_feat = []
175     test_names = []
176     for idx in range(len(file_list)):
177         try:
178             img = cv2.imread(os.path.join(path,file_list[idx]))
179             gram_desc,_ = get_gram_descriptor(img, model, coarse=coarse)
180             gram_feat.append(gram_desc)
181             test_names.append(file_list[idx])
182         except:
183             print(file_list[idx])
184     return np.array(gram_feat), test_names
185
186 def get_class_adain(path,model, class_name="cloudy"):
187     # this is to get the adain features, similar to the previous lbp and
188     # gram matrix features
189     file_list = [x for x in os.listdir(path) if (class_name in x and x.
190     endswith(".jpg"))]
191     adain_feat = []
192     test_names = []
193     for idx in range(len(file_list)):
194         try:
195             img = cv2.imread(os.path.join(path,file_list[idx]))
196             adain_desc = get_adain_features(img, model)
197             adain_feat.append(adain_desc)
198             test_names.append(file_list[idx])
199         except:
200             print(file_list[idx])
201     return np.array(adain_feat), test_names
202
203 def create_dataset(train_data, test_data, type="lbp", downsam=2048):
204     # we build the dataset for each class in this way
205
206     if type=="lbp":
207         # for lbp we just stack the vectors on top of each other
208         train_x = np.vstack((train_data["cloudy"],train_data["rain"],
209         train_data["shine"],train_data["sunrise"]))
210
211         test_x = np.vstack((test_data["cloudy"],test_data["rain"],
212         test_data["shine"],test_data["sunrise"]))
213         # for the labels we just place as many 0s as images in cloudy
214         class, and so on

```



```

208     train_y = np.hstack((np.array([[0]]).repeat(train_data["cloudy"].
209     shape[0]),
210     np.array([[1]]).repeat(train_data["rain"].
211     shape[0]),
212     np.array([[2]]).repeat(train_data["shine"].
213     shape[0]),
214     np.array([[3]]).repeat(train_data["sunrise"].
215     shape[0])))
216     test_y = np.hstack((np.array([[0]]).repeat(test_data["cloudy"].
217     shape[0]),
218     np.array([[1]]).repeat(test_data["rain"].shape
219     [0]),
220     np.array([[2]]).repeat(test_data["shine"].
221     shape[0]),
222     np.array([[3]]).repeat(test_data["sunrise"].
223     shape[0])))
224     elif type=="gram":
225         # for this we flatten the array, and only sample the first 2048
226         values to be used as descriptors
227         train_x = np.vstack((train_data["cloudy"].reshape(train_data["
228         cloudy"].shape[0],train_data["cloudy"].shape[1]*train_data["cloudy"].
229         shape[2])[:, :downsam],
230         train_data["rain"].reshape(train_data["rain"
231         ].shape[0],train_data["rain"].shape[1]*train_data["rain"].shape[2])[:, :
232         downsam],
233         train_data["shine"].reshape(train_data["shine
234         "].shape[0],train_data["shine"].shape[1]*train_data["shine"].shape[2])
235        [:, :downsam],
236         train_data["sunrise"].reshape(train_data["
237         sunrise"].shape[0],train_data["sunrise"].shape[1]*train_data["sunrise"
238         "].shape[2])[:, :downsam]))
239     test_x = np.vstack((test_data["cloudy"].reshape(test_data["cloudy"
240     "].shape[0],test_data["cloudy"].shape[1]*test_data["cloudy"].shape[2])
241    [:, :downsam],
242     test_data["rain"].reshape(test_data["rain"].
243     shape[0],test_data["rain"].shape[1]*test_data["rain"].shape[2])[:, :
244     downsam],
245     test_data["shine"].reshape(test_data["shine"].
246     shape[0],test_data["shine"].shape[1]*test_data["shine"].shape[2])[:, :
247     downsam],
248     test_data["sunrise"].reshape(test_data["
249     sunrise"].shape[0],test_data["sunrise"].shape[1]*test_data["sunrise"].
250     shape[2])[:, :downsam]))
251     train_y = np.hstack((np.array([[0]]).repeat(train_data["cloudy"].
252     shape[0]),
253     np.array([[1]]).repeat(train_data["rain"].
254     shape[0]),
255     np.array([[2]]).repeat(train_data["shine"].
256     shape[0]),
257     np.array([[3]]).repeat(train_data["sunrise"].
258     shape[0])))

```

```

232     test_y = np.hstack((np.array([[0]]).repeat(test_data["cloudy"].
shape[0]),
233                               np.array([[1]]).repeat(test_data["rain"].shape
[0]),
234                               np.array([[2]]).repeat(test_data["shine"].
shape[0]),
235                               np.array([[3]]).repeat(test_data["sunrise"].
shape[0])))
236     elif type=="adain":
237         # we just need to stack since its only the mean and std per class
238         train_x = np.vstack((train_data["cloudy"],train_data["rain"],
train_data["shine"],train_data["sunrise"]))
239
240         test_x = np.vstack((test_data["cloudy"],test_data["rain"],
test_data["shine"],test_data["sunrise"]))
241
242         train_y = np.hstack((np.array([[0]]).repeat(train_data["cloudy"].
shape[0]),
243                                   np.array([[1]]).repeat(train_data["rain"].
shape[0]),
244                                   np.array([[2]]).repeat(train_data["shine"].
shape[0]),
245                                   np.array([[3]]).repeat(train_data["sunrise"].
shape[0])))
246         test_y = np.hstack((np.array([[0]]).repeat(test_data["cloudy"].
shape[0]),
247                                   np.array([[1]]).repeat(test_data["rain"].shape
[0]),
248                                   np.array([[2]]).repeat(test_data["shine"].
shape[0]),
249                                   np.array([[3]]).repeat(test_data["sunrise"].
shape[0])))
250     return train_x, train_y, test_x, test_y
251
252 def get_classified(test_y, test_preds, test_names):
253     # I use this to print out correct and incorrect per classifier
254     # we get all the indices where each class is
255     idx_cloudy = np.argwhere(test_y == 0)[: ,0]
256     idx_rain = np.argwhere(test_y == 1)[: ,0]
257     idx_shine = np.argwhere(test_y == 2)[: ,0]
258     idx_sunrise = np.argwhere(test_y == 3)[: ,0]
259     # then we loop through each and check if its correct or not, and save
it to this dictionary
260     # it will only give the final values, but that is ok since we just
want one example
261     cloudy = {}
262     for idx in idx_cloudy:
263         if test_y[idx] == test_preds[idx]:
264             cloudy["correct"] = test_names[idx]
265         else:
266             # we have this = to a list since we want the class it was
classified as
267             cloudy["incorrect"] = [test_names[idx],test_preds[idx]]
268

```

```
269 rain = {}
270 for idx in idx_rain:
271     if test_y[idx] == test_preds[idx]:
272         rain["correct"] = test_names[idx]
273     else:
274         rain["incorrect"] = [test_names[idx], test_preds[idx]]
275
276 shine = {}
277 for idx in idx_shine:
278     if test_y[idx] == test_preds[idx]:
279         shine["correct"] = test_names[idx]
280     else:
281         shine["incorrect"] = [test_names[idx], test_preds[idx]]
282
283 sunrise = {}
284 for idx in idx_sunrise:
285     if test_y[idx] == test_preds[idx]:
286         sunrise["correct"] = test_names[idx]
287     else:
288         sunrise["incorrect"] = [test_names[idx], test_preds[idx]]
289
290 # we calculate the correct number of classifications by summing all
291 the True in this array and dividing by the total amount of
292 classifications
293 correct = (test_y == test_preds).sum()
294 accuracy = correct/len(test_y)
295 # we just print it
296 print("Accuracy: ", accuracy)
297 # print correct incorrect pairs
298 print(cloudy, rain, shine, sunrise)
299
300 def plot_confusion_matrix(svm, test_y, preds, name, classes=["cloudy", "
rain", "shine", "sunrise"]):
301     cm = confusion_matrix(test_y, preds, labels=svm.classes_)
302     plt.cla()
303     plt.clf()
304     disp = ConfusionMatrixDisplay(confusion_matrix=cm,
305                                   display_labels=classes)
306     disp.plot()
307     plt.savefig(name+"_cm.png", bbox_inches='tight')
308
309 def plot_lbp(lbp_descriptor, name):
310     # this plots the lbp histogram
311     vals = np.arange(len(lbp_descriptor))
312     plt.bar(vals, lbp_descriptor, color='blue',
313            width = 0.8)
314
315     plt.xlabel("Encoding")
316     plt.ylabel("Frecuency")
317     plt.title("LBP Histogram")
318     new_name = name + "_lbp.png"
319     plt.savefig(new_name, bbox_inches='tight')
320     plt.clf()
321     plt.cla()
322
323 def plot_gram(gram_matrix, name, model):
```

```
320     # plots the 512 by 512 gram matrix
321     # we add this 1e-10 to prevent having a log(0)
322     # we want log scale since the values could be big
323     # we are using the unnormalized gram matrix
324     gram_matrix += 1e-10
325     gram_matrix = np.log10(gram_matrix)
326     plt.clf()
327     plt.cla()
328     plt.imshow(gram_matrix, cmap="gray")
329     plt.colorbar()
330     plt.title("Gram Matrix")
331     new_name = name + "_gram_" + model + ".png"
332     plt.savefig(new_name, bbox_inches='tight')
333     plt.clf()
334     plt.cla()
335
336 img_path = "data/training"
337 test_path = "data/testing"
338 classes = ["cloudy", "rain", "shine", "sunrise"]
339 class_dict = {"cloudy": 0, "rain": 1, "shine": 2, "sunrise": 3}
340
341 # LBP
342 lbp_cfeat = {}
343 lbp_test = {}
344 P = 12
345 R = 2
346 names = {}
347 for cls in classes:
348     lbp_cfeat[cls], _ = get_class_lbp(img_path, P=P, R=R, class_name=cls)
349     lbp_test[cls], names[cls] = get_class_lbp(test_path, P=P, R=R,
350                                              class_name=cls)
351 test_filenames = names["cloudy"] + names["rain"] + names["shine"] + names[
352     "sunrise"]
353
354 lbp_train_x, lbp_train_y, lbp_test_x, lbp_test_y = create_dataset(
355     lbp_cfeat, lbp_test, type="lbp")
356
357 svm_lbp = svm.SVC()
358 svm_lbp.fit(lbp_train_x, lbp_train_y);
359
360 lbp_preds = svm_lbp.predict(lbp_test_x)
361
362 get_classified(lbp_test_y, lbp_preds, test_filenames)
363
364 plot_confusion_matrix(svm_lbp, lbp_test_y, lbp_preds, "lbp");
365
366 # GRAM MATRIX
367 vgg = VGG19()
368 vgg.load_weights('vgg_normalized.pth')
369 encoder_name='resnet50'
370 resnet = CustomResNet(encoder=encoder_name)
371
372 gram_cfeat_resnet_coarse = {}
373 gram_test_resnet_coarse = {}
```

```
371 names_resnetc = {}
372
373 gram_cfeat_resnet_fine = {}
374 gram_test_resnet_fine = {}
375 names_resnetf = {}
376
377 gram_cfeat_vgg = {}
378 gram_test_vgg = {}
379 names_vgg = {}
380
381 for cls in classes:
382     gram_cfeat_resnet_coarse[cls], _ = get_class_gram(img_path, resnet,
383     class_name=cls, coarse=True)
384     gram_test_resnet_coarse[cls], names_resnetc[cls] = get_class_gram(
385     test_path, resnet, class_name=cls, coarse=True)
386
387     gram_cfeat_resnet_fine[cls], _ = get_class_gram(img_path, resnet,
388     class_name=cls, coarse=False)
389     gram_test_resnet_fine[cls], names_resnetf[cls] = get_class_gram(
390     test_path, resnet, class_name=cls, coarse=False)
391
392     gram_cfeat_vgg[cls], _ = get_class_gram(img_path, vgg, class_name=cls,
393     coarse=None)
394     gram_test_vgg[cls], names_vgg[cls] = get_class_gram(test_path, vgg,
395     class_name=cls, coarse=None)
396
397 resnetc_train_x, resnetc_train_y, resnetc_test_x, resnetc_test_y =
398     create_dataset(gram_cfeat_resnet_coarse, gram_test_resnet_coarse, type=
399     "gram")
400 resnetf_train_x, resnetf_train_y, resnetf_test_x, resnetf_test_y =
401     create_dataset(gram_cfeat_resnet_fine, gram_test_resnet_fine, type="
402     gram")
403 vgg_train_x, vgg_train_y, vgg_test_x, vgg_test_y = create_dataset(
404     gram_cfeat_vgg, gram_test_vgg, type="gram")
405
406 svm_resnetc = svm.SVC()
407 svm_resnetc.fit(resnetc_train_x, resnetc_train_y);
408 resnetc_preds = svm_resnetc.predict(resnetc_test_x)
409
410 get_classified(resnetc_test_y, resnetc_preds, test_filenames)
411
412 plot_confusion_matrix(svm_lbp, lbp_test_y, resnetc_preds, "resnetc");
413
414 svm_resnetf = svm.SVC()
415 svm_resnetf.fit(resnetf_train_x, resnetf_train_y);
416 resnetf_preds = svm_resnetf.predict(resnetf_test_x)
417
418 get_classified(lbp_test_y, resnetf_preds, test_filenames)
419
420 plot_confusion_matrix(svm_lbp, lbp_test_y, resnetf_preds, "resnetf");
421
422 svm_vgg = svm.SVC()
423 svm_vgg.fit(vgg_train_x, vgg_train_y);
424 vgg_preds = svm_vgg.predict(vgg_test_x)
```

```
414 get_classified(lbp_test_y, vgg_preds, test_filenames)
415
416 plot_confusion_matrix(svm_lbp, lbp_test_y, vgg_preds, "vgg");
417
418 adain_train_vgg = {}
419 adain_test_vgg = {}
420 names_vgg = {}
421
422
423 for cls in classes:
424     adain_train_vgg[cls], _ = get_class_adain(img_path, vgg, class_name=cls
425     )
426     adain_test_vgg[cls], names_vgg[cls] = get_class_adain(test_path, vgg,
427     class_name=cls)
428
429 adain_train_x, adain_train_y, adain_test_x, adain_test_y = create_dataset(
430     adain_train_vgg, adain_test_vgg, type="adain")
431
432 svm_adain= svm.SVC()
433 svm_adain.fit(adain_train_x, adain_train_y);
434 adain_preds = svm_adain.predict(adain_test_x)
435
436 get_classified(lbp_test_y, adain_preds, test_filenames)
437
438 plot_confusion_matrix(svm_lbp, lbp_test_y, adain_preds, "adain");
439
440 ex_path = "data/training"
441 examples = ["cloudy1.jpg", "rain1.jpg", "shine1.jpg", "sunrise1.jpg"]
442
443 for example in examples:
444     basename = example[:-4]
445     img = cv2.imread(os.path.join(ex_path, example))
446     # plot lbp
447     lbp_desc = get_lbp_descriptor(img, P=P, R=R)
448     plot_lbp(lbp_desc, basename);
449
450     # plot gram
451     _, gram_matrix = get_gram_descriptor(img, vgg, coarse=None)
452     plot_gram(gram_matrix, basename, "vgg");
453
454     _, gram_matrix = get_gram_descriptor(img, resnet, coarse=False)
455     plot_gram(gram_matrix, basename, "resnet_fine");
456
457     _, gram_matrix = get_gram_descriptor(img, resnet, coarse=True)
458     plot_gram(gram_matrix, basename, "resnet_coarse");
```

Listing 1: Source code