

Homework 8

Theory Question

1. This idea is fundamental because it is the principle from which we are able to estimate the intrinsic parameters. Since we can get the homographies per camera pose using the chessboard pattern, we can also estimate the parameters of the image of the absolute conic, ω , as shown in Eq. 3, and that allows us to build a linear system of equations which we can solve to get these each of the values of the matrix ω .
2. This is algebraically proven by the fact that $\omega = \mathbf{H}^\top \mathbf{I}_{3 \times 3} \mathbf{H}^{-1}$. Expanding this expression:

$$\begin{aligned}
 \omega &= \mathbf{H}^\top \mathbf{I}_{3 \times 3} \mathbf{H}^{-1} \\
 &= (\mathbf{KR})^{-\top} \mathbf{I}_{3 \times 3} (\mathbf{KR})^{-1} \\
 &= ((\mathbf{KR})^\top)^{-1} (\mathbf{KR})^{-1} \\
 &= (\mathbf{R}^\top \mathbf{K}^\top)^{-1} (\mathbf{R}^{-1} \mathbf{K}^{-1}) \\
 &= \mathbf{K}^{-\top} \mathbf{R}^{-\top} (\mathbf{R}^{-1} \mathbf{K}^{-1}) \\
 &= \mathbf{K}^{-\top} (\mathbf{RR}^\top)^{-1} \mathbf{K}^{-1} \\
 &= \mathbf{K}^{-\top} \mathbf{K}^{-1}
 \end{aligned} \tag{1}$$

Since we know that \mathbf{K} is a triangular matrix, so will its transpose. Additionally, any triangular matrix multiplied by its transpose will always result in a positive-definite matrix. So this proves that the image of the absolute conic will not contain any real pixel locations, all of them are imaginary.

Task: Zhang's algorithm

First of all, we define the 3D points for the pattern corners. For the first dataset we define the length of the black squares as 1, and the grid shape is 8×10 . The order of the points is left to right, top to bottom.

Corner detection

Each image is opened as grayscale, which is then passed through Canny algorithm in order to find the edges of the black squares. The edges are then passed through a Hough transform in order to find the lines. The lines are then sorted depending on the parameters θ and ρ to check if they are either a horizontal line or a vertical line. For horizontal lines we have that $\theta \leq \pi/4$ and $\rho \geq 0$ or $\theta \geq 3\pi/4$ and $\rho \leq 0$, and if they are not horizontal then they are vertical. We then sort horizontal and vertical lines depending on their intersection to the y and x axis (the very top and the far left of the image) respectively. We find lines that have intersections that are within an specified distance of each other, we define this as the distance tolerance value, and we use 15 for the first dataset.

Now that the lines are all sorted in their respective groups, we focus on the groups of lines that have more than 1 line. To handle such cases, we get the intersection of each of the lines

in a group with the top and bottom of the image in the case of vertical lines, and with the left and right of the image in the case of horizontal lines. Then we find the midpoint between these intersections for the lines within a group. Finally, we get our final line, one that passes through the midpoints and keep that. This is all done with lines in their homogeneous coordinates representation.

After our filtering and grouping, we loop through the lines and get the x and y intercept, but this time we get the sorted arguments, so then we sort them with this. Now that the lines are sorted, we use a nested loop, with the inside looping through the vertical lines and the outside loop through the horizontal ones. Since these are in homogeneous coordinates, we can easily find the intersection between them, we store each intersection in an array. The index of each of these intersections in the array belong to the label of the corner.

Calculating Homographies

We use the corners of each of the images and we take the 3D points (without the Z values) and get the corresponding homography for each image with point to point correspondences. We then store each homography for later use.

Estimating intrinsic parameters

We have that the image of the absolute conic Ω_∞ is given by

$$\omega = K^\top K^{-1} \quad (2)$$

Any plane in the world frame samples the absolute conic at exactly two points, and plugging these two points into the equation above gives us:

$$\begin{aligned} \mathbf{h}_1 \omega \mathbf{h}_1 &= \mathbf{h}_2 \omega \mathbf{h}_2 \\ \mathbf{h}_1 \omega \mathbf{h}_2 &= 0 \end{aligned} \quad (3)$$

We know that the structure of the matrix ω is:

$$\omega = \begin{pmatrix} \omega_{11} & \omega_{12} & \omega_{13} \\ \omega_{21} & \omega_{22} & \omega_{23} \\ \omega_{31} & \omega_{32} & \omega_{33} \end{pmatrix}$$

But since it is symmetric, there are only 6 parameters we need to find. Expanding the equations in Eq. 3, we arrive at

$$\begin{aligned} \mathbf{v}_{12}^\top \mathbf{b} &= 0 \\ (\mathbf{v}_{11} - \mathbf{v}_{22})^\top \mathbf{b} &= 0 \end{aligned} \quad (4)$$

With $\mathbf{b} = (\omega_{11} \ \omega_{12} \ \omega_{22} \ \omega_{13} \ \omega_{23} \ \omega_{33})^\top$ Where \mathbf{v} is given by:

$$\mathbf{v}_{ij} = \begin{pmatrix} h_{i1}h_{j1} \\ h_{i1}h_{j2} + h_{i2}h_{j1} \\ h_{i2}h_{j2} \\ h_{i2}h_{j3} + h_{i3}h_{j2} \\ h_{i1}h_{j3} + h_{i3}h_{j1} \\ h_{i3}h_{j3} \end{pmatrix}$$

Rewriting Eq. 4, we have:

$$\mathbf{V}\mathbf{b} = \mathbf{0} \quad (5)$$

Where $\mathbf{V} = (\mathbf{v}_{12}^\top \ (\mathbf{v}_{11} - \mathbf{v}_{22})^\top)^\top$. This is a 2×6 matrix of values that we can calculate. Each position of the camera showing the $Z = 0$ plane will give us 2 equations, since we have 6 unknowns we need 3 camera positions at least. We can stack together the all the equations from the camera positions (homographies) and solve using linear least squares.

After getting the elements of $\boldsymbol{\omega}$ we can get the intrinsic parameters. The matrix \mathbf{K} has the following form:

$$\mathbf{K} = \begin{pmatrix} \alpha_x & s & x_0 \\ 0 & \alpha_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \quad (6)$$

Which arrives us at the relationship between these parameters and $\boldsymbol{\omega}$

$$\begin{aligned} y_0 &= \frac{\omega_{12}\omega_{13} - \omega_{11}\omega_{23}}{\omega_{11}\omega_{22} - \omega_{12}^2} \\ \lambda &= \omega_{33} - \frac{\omega_{13}^2 + y_0(\omega_{12}\omega_{13} - \omega_{11}\omega_{23})}{\omega_{11}} \\ \alpha_x &= \sqrt{\frac{\lambda}{\omega_{11}}} \\ \alpha_y &= \sqrt{\frac{\lambda\omega_{11}}{\omega_{11}\omega_{22} - \omega_{12}^2}} \\ s &= -\frac{\omega_{12}\alpha_x^2\alpha_y}{\lambda} \\ x_0 &= \frac{sy_0}{\alpha_y} - \frac{\omega_{13}\alpha_x^2}{\lambda} \end{aligned} \quad (7)$$

So with this our intrinsic matrix is found.

Estimating extrinsic parameters

The relationship between the intrinsic parameters, the homographies and the extrinsic parameters is given by:

$$\mathbf{K}^{-1} (\mathbf{h}_1 \ \mathbf{h}_2 \ \mathbf{h}_3) = (\mathbf{r}_1 \ \mathbf{r}_2 \ \mathbf{t}) \quad (8)$$

So now, we expand and get a number of equations

$$\begin{aligned} \mathbf{r}_1 &= \zeta \mathbf{K}^{-1} \mathbf{h}_1 \\ \mathbf{r}_2 &= \zeta \mathbf{K}^{-1} \mathbf{h}_2 \\ \mathbf{r}_3 &= \mathbf{r}_1 \times \mathbf{r}_2 \\ \mathbf{t} &= \zeta \mathbf{K}^{-1} \mathbf{h}_3 \end{aligned} \quad (9)$$

The variable ζ is a scale factor to enforce that all the columns of the \mathbf{R} matrix are of unit magnitude, with a value of $\zeta = 1 / \|\mathbf{K}^{-1} \mathbf{h}_1\|$.

Now \mathbf{R} and \mathbf{t} are known. However, there is a few more steps we must follow. The rotation matrix obtained is still not guaranteed to be orthonormal. For this, we decompose \mathbf{R} with SVD, resulting in $\mathbf{R} = \mathbf{U}\mathbf{D}\mathbf{V}^\top$ and set $\mathbf{R} = \mathbf{U}\mathbf{V}^\top$.

For our next step we need to have another representation of the rotation matrix, as its DoF is 3, but has 9 values. We want the number of values of this representation to be the same number as the DoF. For this we get the Rodrigues vector, \mathbf{w} , given by:

$$\mathbf{w} = \frac{\phi}{2 \sin \phi} \begin{pmatrix} r_{32} - r_{23} \\ r_{13} - r_{31} \\ r_{21} - r_{12} \end{pmatrix}$$

$$\mathbf{R} = \begin{pmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{pmatrix} \quad (10)$$

And to go from the Rodrigues vector back to the rotation matrix we follow:

$$\mathbf{R} = \mathbf{I}_{3 \times 3} + \frac{\sin \phi}{\phi} [\mathbf{w}]_X + \frac{1 - \cos \phi}{\phi^2} [\mathbf{w}]_X^2$$

$$\phi = \|\mathbf{w}\|$$

$$[\mathbf{w}]_X = \begin{pmatrix} 0 & -w_z & w_y \\ w_z & 0 & -w_x \\ -w_y & w_x & 0 \end{pmatrix} \quad (11)$$

Parameter Refinement

Now that we have all our intrinsic and extrinsic parameters we need to refine these. We do so with Levenberg-Marquadt algorithm. For this we need to optimize the reprojection error, which will be our Euclidean distance. For this we separate all parameters and optimize for it. One thing that we consider is that the intrinsic parameters are shared between all images, so we implement our optimization to reflect that.

Plotting Camera Poses

After refining the parameters with non-linear least squares, we can now find the position of the camera, the camera plane and the direction of the camera's coordinate axis. For that we make use of the extrinsic parameters in the following manner

$$\mathbf{C} = -\mathbf{R}^\top \mathbf{t}$$

$$\mathbf{X}_{cam} = -\mathbf{R}\mathbf{X} + \mathbf{t} \quad (12)$$

$$\mathbf{X} = \mathbf{R}^{-1} \mathbf{X}_{cam} - \mathbf{R}^{-1} \mathbf{t}$$


$$= \mathbf{R}^\top \mathbf{X}_{cam} - \mathbf{R}^\top \mathbf{t} = \mathbf{R}^\top \mathbf{X}_{cam} + \mathbf{C} \quad (13)$$

$$\mathbf{X}_{cam}^x = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \mathbf{X}_{cam}^y = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \mathbf{X}_{cam}^z = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} \quad (14)$$


With this we now have the camera plane (the plane that contains \mathbf{X}_{cam}^x and \mathbf{X}_{cam}^y). One modification we do for the purpose of plotting is that we don't consider adding \mathbf{C} in Eq. 13, as the library we are using to make the 3D plot takes in the origin of the vector and the direction, so there is no need here to add \mathbf{C} back in.

Results: Dataset 1


We will be using 'Pic_19.jpg' and 'Pic_31.jpg' to show the corners, the reprojection, the comparison between the refined parameters and the unrefined ones, and 'Pic_11.jpg' is our 'fixed image'.



(a) 'Pic_11.jpg', our fixed image



(b) 'Pic_19.jpg'.



(c) 'Pic_31.jpg'.

Figure 1: Example poses.


Corner detection

As mentioned before, the process of detecting corners is:


1. Use Canny algorithm on the grayscale image. We make use of OpenCV's Canny function, with thresholds of 300 and 350.

2. Use Hough transform on the Canny edges. Again, we use OpenCV's HoughLines function, with a parameter of 45.
3. We separate lines into vertical and horizontal, group those which are close and find the "midpoint" line for each group.
4. Now we sort all the vertical and horizontal lines depending on their x and y intercept.
5. Get the intersection of each vertical and horizontal line.


The process can be seen for the fixed image and our other 2 images in Figs. 2, 3 and 4.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.




(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.




(d) Detected corners.


Figure 2: Corner detection process with 'Pic_11.jpg'.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.




(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.




(d) Detected corners.


Figure 3: Corner detection process with 'Pic_19.jpg'.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.



(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.



(d) Detected corners.

Figure 4: Corner detection process with 'Pic_31.jpg'.

Intrinsic Parameters

First we found the homographies for each image through the use of point to point correspondences with the point grid without considering the z coordinates and store them in a list. We then loop through this list and get the equations shown in Eq. 4. We stack all the matrices \mathbf{V} found per image and solve it using Linear Least Squares. So now we get \mathbf{K} using the ω found. Our initial values found for \mathbf{K} are:

$$\mathbf{K} = \begin{pmatrix} 717.9290664 & 0.80913725 & 236.27489463 \\ 0 & 715.18947792 & 321.37701144 \\ 0 & 0 & 1 \end{pmatrix}$$

Extrinsic parameters

We use the intrinsic parameters to calculate the rotation matrix and the translation vector according to the equations described in the extrinsic parameter section, we also condition

the rotation matrix and then append both the rotation matrix and the translation vector to an array for further use. Here are two examples, this is the extrinsics for 'Pic_19.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.85453588 & 0.08678829 & 0.51209006 & -2.38912397 \\ -0.04449309 & 0.99454832 & -0.09430805 & -6.58271785 \\ -0.51748314 & 0.05780514 & 0.8537387 & 19.28381236 \end{pmatrix}$$

And for 'Pic_31.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.9795844 & 0.15225277 & 0.13127638 & -1.93246888 \\ -0.14268342 & 0.98656977 & -0.07950799 & -5.77755365 \\ -0.14161862 & 0.05915382 & 0.98815231 & 18.37295395 \end{pmatrix}$$

Optimizing the parameters

To refine our parameters, we place the intrinsic parameters $\alpha_x, s, x_0, \alpha_y, y_0$ in an array. Then we loop through the lists containing the rotation matrices and the translation vectors. We obtain the Rodrigues vector from each of the rotation matrices, and append the 3 parameters of this vector to the parameter array, we do the same with the 3 parameters of the translation vector. After having all the parameters in an array, we can finally call Scipy's Levenberg-Marquadt algorithm. The function we will optimize is a cost function that we define to take in the the parameters, the 3D coordinates and the corners obtained from the corner detection process. This function will use the parameters to obtain the \mathbf{K} , \mathbf{R} and \mathbf{t} then use these to obtain the projection matrix given by $\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}]$ per image.

To visually check the improvement, we used 'Pic_11.jpg' as our 'Fixed Image'. We projected the corners for both image 'Pic_19.jpg' and 'Pic_31.jpg' to the 3D space (but only on the $Z = 0$ plane), then we took these points and projected them into the fixed image using its projection matrix (only the first, second and last column were used).

After our optimization our \mathbf{K} matrix looks like:

$$\mathbf{K} = \begin{pmatrix} 811.00158749 & 5.56088429 & 263.17953178 \\ 0 & 813.5051964 & 229.06999829 \\ 0 & 0 & 1 \end{pmatrix}$$

And now, the optimized extrinsics for 'Pic_19.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.84324895 & 0.10543803 & 0.52708067 & -2.97020936 \\ -0.0745388 & 0.9940363 & -0.0795978 & -3.93242111 \\ -0.53232996 & 0.0278328 & 0.84607928 & 20.94152782 \end{pmatrix}$$

And for 'Pic_31.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.97823119 & 0.15597661 & 0.136876 & -2.58454503 \\ -0.14537129 & 0.98577176 & -0.08438737 & -3.36038005 \\ -0.14809095 & 0.06265252 & 0.9869872 & 20.52779866 \end{pmatrix}$$

We can see that indeed, LM optimization significantly reduced the error. In fact, it is so accurate that it is hard to tell where the reprojected corners are in some of the cases. This

Image	Error mean	Error mean (after LM)	Error std	Error std (after LM)
'Pic_19.jpg'	10.36677	1.086268	4.89901	0.496179
'Pic_31.jpg'	2.620455	0.870828	1.08557	0.4628802

Table 1: Euclidean error mean and standard deviation for the reprojection error into the fixed image.


is shown visually in Figs. 5 and 6, the reprojected corners onto the fixed image are vastly better after LM. Another metrics we can use to gauge this error is by using the mean and the standard deviation, and as we can see in Table 1, LM algorithm massively improved both, lowering the error by several times and the standard deviation as well. Also just projecting the 3D coordinates to the images shows improvement, as seen in Figs. 7 and 8.




(a) Fixed image with reprojected corners before LM.

(b) Fixed image with reprojected corners after LM.

Figure 5: Visual reprojection from 'Pic_19.jpg' into the fixed image. The green dots are for the position of the detected corners and red is for the reprojected corners.




(a) Fixed image with reprojected corners before LM.




(b) Fixed image with reprojected corners after LM.

Figure 6: Visual reprojection from 'Pic_31.jpg' into the fixed image. The green dots are for the position of the detected corners and red is for the reprojected corners.

Now we will also try projecting the 3D points into the images using the projection matrix before and after LM optimization, and visually compare with the detected points.




(a) Reprojected corners before LM.




(b) Reprojected corners after LM.

Figure 7: Visual reprojection of the 3D coordinates into 'Pic_19.jpg'. Green is the detected corners and red is the projected corners.



(a) Reprojected corners before LM.



(b) Reprojected corners after LM.

Figure 8: Visual reprojection of the 3D coordinates into 'Pic_31.jpg'. Green is the detected corners and red is the projected corners.

Camera poses

We now take our optimized \mathbf{R} and \mathbf{t} to get the camera poses. We follow the equations and obtain both the camera coordinates and the direction of the camera axis. To plot them we use Matplotlib and art3d from `mpl_toolkits.mplot3d`. We also make sure to plot the calibration pattern in this 3d plot using `Rectangle` from `matplotlib.patches`. To plot this pattern, we notice that it "skips" a line, there is one row with 4 black squares, then the next row is blank, and then 4 black squares again. So we make sure to have this reflected in the plotted calibration pattern. Notice that all the camera centers are in the negative side of the z axis, this is due to when we define the coordinates of the calibration pattern, we set the origin to be at the top left side, and increment the x coordinate when we go right. Similarly, we increase the y coordinate when we go to the bottom of the image, so the z axis is actually pointing towards the image, thus all the camera poses have negative z values. We also make sure to plot the camera plane as a rectangle, to do this we know that this plane has 2 orthonormal vectors, in our case it is the x and y direction vectors of the camera, which we already have. So to plot this rectangle, we need each of the corners. Each of this corners, c_i will be given by:

$$\begin{aligned}
 c_1 &= \mathbf{C} + \mathbf{X}_x + \mathbf{X}_y \\
 c_2 &= \mathbf{C} + \mathbf{X}_x - \mathbf{X}_y \\
 c_3 &= \mathbf{C} - \mathbf{X}_x - \mathbf{X}_y \\
 c_4 &= \mathbf{C} - \mathbf{X}_x + \mathbf{X}_y
 \end{aligned} \tag{15}$$

Where \mathbf{X}_x and \mathbf{X}_y are the directions of the camera x and y axis respectively. We also make sure to randomize the color of the camera plane for each different image.






Figure 9: Camera poses for all the dataset 1 images.

Results: Own dataset


We largely follow the same procedure for the dataset we created. With the differences described in each subsection.



(a) 'fixed_img.jpg', our fixed image for our dataset



(b) 'pose9.jpg'.




(c) 'pose33.jpg'.


Figure 10: Example poses for our own dataset.

Corner detection


We follow the same procedure, except due to the lighting conditions during the time we took the pictures, there is an additional step when filtering the lines. If the detected lines are more than the expected lines, we just take the first few, this can be seen in Figs. [11](#), [12](#), [13](#). The reason why we can do this is because the extra line are in the far right, there was also a case with a rogue line in the far bottom, but we clean it up in the same way. We use the same parameters for Canny and HoughLines as we did for dataset 1.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.




(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.




(d) Detected corners.


Figure 11: Corner detection process with 'fixed_img.jpg'.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.




(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.




(d) Detected corners.


Figure 12: Corner detection process with 'pose9.jpg'.




(a) Edges from Canny algorithm.



(b) Detected lines with Hough line transform.



(c) Filtered and 'averaged' lines, vertical lines are shown in blue and horizontal are shown in green.



(d) Detected corners.

Figure 13: Corner detection process with 'pose33.jpg'.

Intrinsic Parameters

The matrix \mathbf{K} found is:

$$\mathbf{K} = \begin{pmatrix} 470.50104806 & -0.5421189 & 240.79357019 \\ 0 & 470.34740048 & 317.93433605 \\ 0 & 0 & 1 \end{pmatrix}$$

Extrinsic Parameters

We followed the same procedure as for dataset 1. This is the extrinsics matrix for our 'fixed_img.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.9983198 & 0.00993862 & -0.05708583 & -1.85545655 \\ -0.01205899 & 0.99924547 & -0.03691993 & -6.36269796 \\ 0.05667582 & 0.0375463 & 0.99768639 & 12.95199808 \end{pmatrix}$$

Since we are using a length of 1 inch for the black squares, the very last value of the translation vector means the pattern was 12.95 inches, which is close to the actual distance. We took this image approximately 13.2 inches away from the pattern. And for 'pose9.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.99403331 & 0.00932924 & -0.10867728 & -0.9852933 \\ -0.04960155 & 0.92602253 & -0.37419507 & -4.99084077 \\ 0.09714665 & 0.37735293 & 0.92095999 & 10.16441946 \end{pmatrix}$$

The extrinsics for 'pose33.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.89904183 & 0.02200944 & 0.43730923 & -1.7860764 \\ 0.04535019 & 0.98868418 & -0.14299286 & -7.15417657 \\ -0.43550791 & 0.14838862 & 0.8878703 & 14.6940454 \end{pmatrix}$$

Optimizing the parameters

Our optimized \mathbf{K} is:

$$\mathbf{K} = \begin{pmatrix} 469.36153145 & -0.02628371 & 319.06039838 \\ 0 & 469.81047284 & 240.74777318 \\ 0 & 0 & 1 \end{pmatrix}$$

The extrinsics for 'fixed_img.jpg' are:

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.99829642 & 0.00181529 & -0.05831785 & -4.03596233 \\ -0.00386343 & 0.99937891 & -0.03502673 & -4.23588763 \\ 0.05821805 & 0.03519237 & 0.99768339 & 12.98397892 \end{pmatrix}$$

We can see that the last value of \mathbf{t} is now closer to the measured value of 13.2. And for 'pose9.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.99403331 & 0.00932924 & -0.10867728 & -2.71841043 \\ -0.04960155 & 0.92602253 & -0.37419507 & -3.37535341 \\ 0.09714665 & 0.37735293 & 0.92095999 & 10.30325731 \end{pmatrix}$$

The extrinsics for 'pose33.jpg':

$$[\mathbf{R}|\mathbf{t}] = \begin{pmatrix} 0.9172927 & 0.03538065 & 0.39663876 & -3.99275192 \\ 0.01674614 & 0.9917367 & -0.12719232 & -4.49062442 \\ -0.39786136 & 0.12331476 & 0.90912035 & 13.85745156 \end{pmatrix}$$

Now, followed the same procedure of reprojecting corners into the fixed image, the results for the mean and the standard deviation are shown in Table 2.


We can visually see that the reprojection got significantly better after the optimization.

Image	Error mean	Error mean (after LM)	Error std	Error std (after LM)
'pose33.jpg'	13.92340	1.346285	6.142546	0.703136
'pose9.jpg'	9.113768	1.042369	6.007312	0.51973

Table 2: Euclidean error mean and standard deviation for the reprojection error into the fixed image.



(a) Fixed image with reprojected corners before LM.



(b) Fixed image with reprojected corners after LM.

Figure 14: Visual reprojection from 'pose9.jpg' into the fixed image. The green dots are for the position of the detected corners and red is for the reprojected corners.




(a) Fixed image with reprojected corners before LM.




(b) Fixed image with reprojected corners after LM.

Figure 15: Visual reprojection from 'pose33.jpg' into the fixed image. The green dots are for the position of the detected corners and red is for the reprojected corners.

Now we will also try projecting the 3D points into the images using the projection matrix before and after LM optimization, and visually compare with the detected points.




(a) Reprojected corners before LM.




(b) Reprojected corners after LM.

Figure 16: Visual reprojection of the 3D coordinates into 'pose9.jpg'. Green is the detected corners and red is the projected corners.



(a) Reprojected corners before LM.



(b) Reprojected corners after LM.

Figure 17: Visual reprojection of the 3D coordinates into 'pose33.jpg'. Green is the detected corners and red is the projected corners.

Camera poses

We followed the same procedure, and plotted the camera poses in the same way.




Figure 18: Camera poses for all the dataset 2 images.

Source code

```

1  import cv2
2  import os
3  import numpy as np
4  import matplotlib.pyplot as plt
5  from scipy.optimize import least_squares
6
7
8  # we use these to plot the camera poses
9  from mpl_toolkits.mplot3d.art3d import *
10 from matplotlib.patches import Rectangle
11
12 def get_point_or_line(point1, point2):

```

```

13     # first we check if these are already in homogeneous form or not, if
not we put them in homogeneous representation
14     if len(point1) == 2:
15         point1 = np.array(point1)
16         point1 = np.append(point1, 1)
17     if len(point2) == 2:
18         point2 = np.array(point2)
19         point2 = np.append(point2, 1)
20     # then get the cross product
21     # this function has uses for us as both getting the line
representation and
22     # getting line intersections, as the math is the same
23     r = np.cross(point1, point2)
24     if r[2] != 0:
25         return r/r[2]
26     return r
27
28 def plot_hough_lines(img, lines, name):
29     # this is just a function we use to plot the results of the hough line
detector, prior to filtering and averaging
30     img_ = img.copy()
31     if lines is not None:
32         for i in range(0, len(lines)):
33             rho = lines[i][0][0]
34             theta = lines[i][0][1]
35             a = np.cos(theta)
36             b = np.sin(theta)
37             x0 = a * rho
38             y0 = b * rho
39             pt1 = (int(x0 + 1000*(-b)), int(y0 + 1000*(a)))
40             pt2 = (int(x0 - 1000*(-b)), int(y0 - 1000*(a)))
41             img_ = cv2.line(img_, pt1, pt2, (0,0,255), 1, cv2.LINE_AA)
42     cv2.imwrite(name, img_)
43
44 def plot_line_rep(line, image, horizontal=False):
45     # this so we get the endpoints for the lines to plot after separating
into vertical and horizontal
46     # notice how we calculate different endpoints depending on if its
vertical or not
47     # this is because if its horizontal we want the intercepts with the
left and rightmost part of the image,
48     # and if its vertical we want the same but with top and bottom
49     if horizontal:
50         ref_1 = get_point_or_line([0,10],[0,15])
51         ref_2 = get_point_or_line([image.shape[1]-1,10],[image.shape
[1]-1,15])
52     else:
53         ref_1 = get_point_or_line([10,0],[15,0])
54         ref_2 = get_point_or_line([10,image.shape[0]-1],[15,image.shape
[0]-1])
55     point_1 = get_point_or_line(line, ref_1)[:2]
56     point_2 = get_point_or_line(line, ref_2)[:2]
57     return point_1, point_2
58

```

```

59 def cvl_to_homogeneous(cv2_line):
60     # we use this to convert the line from rho theta to homogeneous
    coordinates
61     # we operate most of the time with homogeneous as it is easier to use
62     rho = cv2_line[0][0]
63     theta = cv2_line[0][1]
64     a = np.cos(theta)
65     b = np.sin(theta)
66     x0 = a * rho
67     y0 = b * rho
68     pt1 = (int(x0 + 1000*(-b)), int(y0 + 1000*(a)))
69     pt2 = (int(x0 - 1000*(-b)), int(y0 - 1000*(a)))
70
71     return get_point_or_line(pt1, pt2)
72
73 def separate_lines(lines):
74     # we separate lines by horizontal lines (theta < pi/4 and rho positive
    , or theta>3pi/4 and negative rho),
75     # and if they are not horizontal then they are vertical
76     horizontal = []
77     vertical = []
78     for line in lines:
79         rho = line[0][0]
80         theta = line[0][1]
81         if (theta <= np.pi/4 and rho >= 0) or (theta > 3*np.pi/4 and rho <
    0):
82             vertical.append(line)
83         else:
84             horizontal.append(line)
85     return vertical, horizontal
86
87 def group_lines(lines, tol=30, image_shape=(600,600), horizontal=False,
    grid_shape=(8,10)):
88     # in here we group lines together with depending on the spot they
    intersect the edges of the image
89     # for horizontal lines we check where it hits the left edge of the
    image (y=0)
90     # for vertical lines we check where it hits the top of the image (x=0)
91     # we group them by checking which intersects are within a distance
    threshold of another intersect
92     # after we have grouped them together, we
93     if horizontal:
94         ref_line = get_point_or_line([0,10],[0,15])
95         ref_line_n = get_point_or_line([image_shape[1],10],[image_shape
    [1],15])
96
97     else:
98         ref_line = get_point_or_line([10,0],[15,0])
99         ref_line_n = get_point_or_line([10,image_shape[0]],[15,image_shape
    [0]])
100     tot = []
101     added = [False] * len(lines)
102     for idx,iline in enumerate(lines):
103         aux = []

```

```

104     iline_rep = cvl_to_homogeneous(iline)
105     ref_inter = get_point_or_line(iline_rep, ref_line)
106     if not added[idx]:
107         aux.append(iline_rep)
108         added[idx] = True
109         for jdx, jline in enumerate(lines):
110             if (jline != iline).any():
111                 jline_rep = cvl_to_homogeneous(jline)
112                 p_inter = get_point_or_line(jline_rep, ref_line)
113                 if horizontal:
114                     dist = np.abs(ref_inter[1] - p_inter[1])
115                 else:
116                     dist = np.abs(ref_inter[0] - p_inter[0])
117                 if dist < tol and added[jdx] == False:
118                     aux.append(jline_rep)
119                     added[jdx] = True
120         tot.append(aux)
121     # now that we have them separated by groups, we "average" those which
122     # have 2 or more lines
123     a_lines = []
124     for lines in tot:
125         if len(lines) > 1:
126             inter_point1 = np.array([0,0], dtype=float)
127             inter_point2 = np.array([0,0], dtype=float)
128             count = 0
129             for line in lines:
130                 inter_point1 += get_point_or_line(ref_line, line)[:2]
131                 inter_point2 += get_point_or_line(ref_line, line)[:2]
132                 count += 1
133             inter_point1 = inter_point1/count
134             inter_point2 = inter_point2/count
135             n_line = get_point_or_line(inter_point1, inter_point2)
136         else:
137             n_line = lines[0]
138         a_lines.append(n_line)
139
140     # now that we have grouped and averaged the lines, we will order them
141     # with their x or y intercept depending on the orientation
142     n_intercept = []
143     for line in a_lines:
144         intercept = get_point_or_line(line, ref_line)
145         if horizontal:
146             n_intercept.append(intercept[1])
147         else:
148             n_intercept.append(intercept[0])
149     # we sort them with their intercepts
150     n_intercept = np.array(n_intercept)
151     inter_idx_sorted = np.argsort(n_intercept)
152     t_lines = []
153     for idx in inter_idx_sorted:
154         t_lines.append(a_lines[idx])

```

```

155     # this is the very last step, it will only filter the excess ones,
156     this is mainly used for our dataset 2
157     if horizontal:
158         if len(t_lines) > grid_shape[1]:
159             t_lines = t_lines[:grid_shape[1]]
160     else:
161         if len(t_lines) > grid_shape[0]:
162             t_lines = t_lines[:grid_shape[0]]
163     return t_lines
164
165 def get_corners(image, tol=15, save_canny=None, save_lines=None,
166               save_hough=None, canny_p1=370, canny_p2=300, hough_p=50):
167     # this function runs canny, houghlines, groups the lines, plots all of
168     this,
169     # also will loop through all the lines and get the intersects, those
170     are our corners.
171     img = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
172     edges = cv2.Canny(img, canny_p1, canny_p2)
173     lines = cv2.HoughLines(edges, 1, np.pi / 180, hough_p, None, 0, 0)
174     if save_hough != None:
175         plot_hough_lines(image, lines, save_hough[:-4]+"_houghlines.jpg")
176     if save_canny != None:
177         cv2.imwrite(save_canny[:-4]+"_canny.jpg", edges)
178
179     vert_lines, hor_lines = separate_lines(lines)
180     v_lines = group_lines(vert_lines, tol=tol, image_shape=(image.shape[1],
181 image.shape[0]), horizontal=False)
182     h_lines = group_lines(hor_lines, tol=tol, image_shape=(image.shape[1],
183 image.shape[0]), horizontal=True)
184     if save_lines != None:
185         line_plot = image.copy()
186         for vline in v_lines:
187             p1, p2 = plot_line_rep(vline, line_plot, horizontal=False)
188             p1_ = (int(p1[0]), int(p1[1]))
189             p2_ = (int(p2[0]), int(p2[1]))
190             line_plot = cv2.line(line_plot, p1_, p2_, (255,0,0), 2);
191         for hline in h_lines:
192             p1, p2 = plot_line_rep(hline, line_plot, horizontal=True)
193             p1_ = (int(p1[0]), int(p1[1]))
194             p2_ = (int(p2[0]), int(p2[1]))
195             line_plot = cv2.line(line_plot, p1_, p2_, (0,255,0), 2);
196         cv2.imwrite(save_lines[:-4]+"_fillines.jpg", line_plot)
197
198     corners = []
199     #this is how we get the corners
200     for h_line in h_lines:
201         for v_line in v_lines:
202             corner = get_point_or_line(h_line, v_line)[:2]
203             corners.append(corner)
204     return corners
205
206 # reusing these from my hw5
207 def point_to_point_system(domain, range_, num_points):
208     # we changed this function to accept an arbitrary number of points,

```

```

203     # this is because most of the systems we will solve have more than 4
data points
204     mat1 = np.empty((0,8),dtype=float)
205     mat2 = np.empty((0,0),dtype=float)
206     for i in range(num_points):
207         x = domain[i,0]
208         x_prime = range_[i,0]
209         y = domain[i,1]
210         y_prime = range_[i,1]
211
212         mat1 = np.append(mat1,np.array([[x, y, 1, 0, 0, 0, -x*x_prime, -y*
x_prime]]),axis=0)
213         mat1 = np.append(mat1,np.array([[0, 0, 0, x, y, 1, -x*y_prime, -y*
y_prime]]),axis=0)
214         mat2 = np.append(mat2,x_prime)
215         mat2 = np.append(mat2,y_prime)
216
217     return mat1,mat2.reshape(num_points*2,1)
218
219 def linear_least_squares(mat1, mat2):
220     # we make a distinction because this is for inhomogeneous least
squares
221     return np.linalg.inv(mat1.T @ mat1) @ mat1.T @ mat2
222
223 def linear_least_squares_homogeneous(mat1):
224     # and this is for homogeneous least squares
225     _, _, v_t = np.linalg.svd(mat1.T @ mat1)
226     return v_t[-1]
227
228 def get_H_matrix(domain,range_):
229     # I reuse this from hw5
230     n_points = len(domain)
231     mat1, mat2 = point_to_point_system(domain, range_,n_points)
232     # we use this function to solve the equation described in the logic,
233     # I replaced np.dot with @ as that is what they suggest in the numpy
website
234     # in this case we use the pseudo inverse (ATA)-1 AT
235     sol = linear_least_squares(mat1,mat2)
236     # we append the 1 since this will only have 8 values, the 1 is missing
237     sol = np.append(sol,np.array([[1]]),axis=0)
238     return sol.reshape((3,3)).astype(float)
239
240 def add_ones(array):
241     # this function takes in an array and adds ones
242     # this is mainly for applying homographies
243     ones_np = np.ones(array.shape[:-1])[...,None]
244     # need to expand the dimensions by 1 to be able to concatenate it
245     return np.append(array,ones_np,axis=-1)
246 def add_number(array,n):
247     # this function takes in an array and adds any number
248     # this is mainly for adding a 0, for z=0 in the real life plane
coordinates
249     ones_np = np.ones(array.shape[:-1])[...,None]
250     ones_np *= n

```



```

251     # need to expand the dimensions by 1 to be able to concatenate it
252     return np.append(array, ones_np, axis=-1)
253
254 def apply_homography(positions, H):
255     # reused from hw5
256     # this gets the homography transformation for all the coordinates in
257     # the image that we get from the get_positions function
258     # we do it in a way that exploits broadcasting, so we don't need to
259     # use for loops
260     temp_pos = add_ones(positions)
261     new_pos = (H @ temp_pos.T).astype(float)
262     new_pos /= new_pos[2,:]
263     return new_pos[:2,:].T
264
265 def get_coordinates(shape=(8,10),length=5):
266     # this is how we create our 2d grid that is on Z=0,
267     # for both datasets we end up using length =1 since the squares are of
268     # 1 inch length
269     x = np.linspace(0, length*(shape[0] - 1), shape[0])
270     y = np.linspace(0, length*(shape[1] - 1), shape[1])
271     xv, yv = np.meshgrid(x,y)
272     xv = xv.ravel()[:, None]
273     yv = yv.ravel()[:, None]
274     return np.hstack((xv,yv))
275
276 def get_vij(H, i,j):
277     # we change ij to go from 0 to 2 rather than 1 to 3, to adjust for
278     # python indexes
279     # we follow the equations from the scrolls
280     v_ij = np.array([H[0,i]*H[0,j],
281                     H[0,i]*H[1,j]+H[0,j]*H[1,i],
282                     H[1,i]*H[1,j],
283                     H[0,i]*H[2,j]+H[0,j]*H[2,i],
284                     H[1,i]*H[2,j]+H[1,j]*H[2,i],
285                     H[2,i]*H[2,j]])
286     return v_ij
287
288 def get_V(H):
289     # this is to get the V matrix
290     v_11 = get_vij(H,0,0)[..., None]
291     v_12 = get_vij(H,0,1)[..., None]
292     v_22 = get_vij(H,1,1)[..., None]
293     V = np.vstack((v_12.T,(v_11 - v_22).T))
294     return V
295
296 def get_b(V):
297     # we get b from the stack of Vs
298     b = linear_least_squares_homogeneous(V)
299     return b
300
301 def get_omega(V):
302     # this is the entire stack of Vs,
303     # we get omega from b which we get from the stack of Vs
304     b = get_b(V)

```

```

301     omega = np.array([[b[0], b[1], b[3]],
302                       [b[1], b[2], b[4]],
303                       [b[3], b[4], b[5]]])
304     return omega
305
306 def get_K_matrix(omega):
307     # we follow the instructions from the scrolls and build our K
308
309     x_0 = (omega[0,1]*omega[0,2] - omega[0,0]*omega[1,2])/(omega[0,0]*
310 omega[1,1] - omega[0,1]**2)
311     lambda_ = omega[2,2] - (omega[0,2]**2 + x_0*(omega[0,1]*omega[0,2] -
312 omega[0,0]*omega[1,2]))/(omega[0,0])
313     alpha_x = np.sqrt(lambda_/omega[0,0])
314     alpha_y = np.sqrt(lambda_*omega[0,0]/(omega[0,0]*omega[1,1] - omega
315 [0,1]**2))
316     s = -1 * (omega[0,1] * alpha_x**2 * alpha_y)/(lambda_)
317     y_0 = (s * x_0 / alpha_y) - (omega[0,2]*alpha_x**2)/(lambda_)
318     K = np.array([[alpha_x, s, x_0],
319                   [0, alpha_y, y_0],
320                   [0, 0, 1]])
321     return K
322
323 def get_extrinsics(K, H):
324     # just following the instructions from the scrolls
325     scale = 1 / np.linalg.norm(np.linalg.inv(K) @ H[:,0])
326     r_1 = scale * np.linalg.inv(K) @ H[:,0]
327     r_2 = scale * np.linalg.inv(K) @ H[:,1]
328     r_3 = np.cross(r_1, r_2)
329     t = scale * np.linalg.inv(K) @ H[:,2]
330     R = np.column_stack((r_1,r_2,r_3))
331     # now we condition R, very important step
332     u, _, v = np.linalg.svd(R)
333     R_conditioned = u @ v
334     return R_conditioned, t
335
336 def get_rodrigues_rep_from_R(R):
337     # we get the rodrigues vector from the rotation matrix, again
338     following the scrolls
339     phi = np.arccos((np.trace(R)-1)/2)
340     w = np.array([R[2,1] - R[1,2],
341                   R[0,2] - R[2,0],
342                   R[1,0] - R[0,1]])
343
344     return w * (phi/(2*np.sin(phi)))
345
346 def get_R_from_rodrigues(w_0, w_1, w_2):
347     # this function is to help us reconstruct our matrices during LM
348     optimization,
349     # more on that in the following functions
350     # again we follow the scrolls
351     W_X = np.array([[0, -w_2, w_1],
352                     [w_2, 0, -w_0],
353                     [-w_1, w_0, 0]])
354     phi = np.sqrt(w_0**2 + w_1**2 + w_2**2)

```

```

350
351     R = np.eye(3) + (np.sin(phi)/phi) * W_X + ((1 - np.cos(phi))/(phi**2))
      * W_X @ W_X
352     return R
353
354 def get_params_from_matrices(K, R_list, t_list):
355     params = []
356     # get the parameters needed from K, R, and t
357     # R and t are per image,
358     # since K is shared between all images, we make sure to optimize it in
      this way, that is why it is only appended once
359     params.append(K[0,0])
360     params.append(K[0,1])
361     params.append(K[0,2])
362     params.append(K[1,1])
363     params.append(K[1,2])
364     for idx in range(len(R_list)):
365         R = R_list[idx]
366         t = t_list[idx]
367         #we need to use rodrigues instead of the rotation matrix
368         w = get_rodrigues_rep_from_R(R)
369         params.append(w[0])
370         params.append(w[1])
371         params.append(w[2])
372         params.append(t[0])
373         params.append(t[1])
374         params.append(t[2])
375     return np.array(params)
376
377 def get_matrices_from_params(params):
378     # this is how we reconstruct our matrices using the parameters
379     alpha_x = params[0]
380     alpha_y = params[1]
381     s = params[2]
382     x_0 = params[3]
383     y_0 = params[4]
384     new_R_list = []
385     new_t_list = []
386     for idx in range(5, len(params), 6):
387         w_0 = params[idx]
388         w_1 = params[idx + 1]
389         w_2 = params[idx + 2]
390         # we get our R back
391         R_ = get_R_from_rodrigues(w_0, w_1, w_2)
392
393         t_0 = params[idx + 3]
394         t_1 = params[idx + 4]
395         t_2 = params[idx + 5]
396         t_ = np.array([t_0, t_1, t_2])
397
398         new_R_list.append(R_)
399         new_t_list.append(t_)
400
401     K = np.array([[alpha_x, s, x_0],

```

```

402         [0, alpha_y, y_0],
403         [0, 0, 1]])
404
405     return K, new_R_list, new_t_list
406
407 def get_projection_matrix(K, R, t):
408     # this is just following the scrolls, this is the definition of
409     # projection matrix
410     R_t = np.hstack((R, t[:,None]))
411     P = K @ R_t
412     return P
413
414 def error_f(params, pts_3d, pts_img_list):
415     # this is the function we use to optimize,
416     # first we get every parameter in a numpy array
417     K, R_list, t_list = get_matrices_from_params(params)
418     error = np.empty((0), dtype=np.float32)
419     # get the error X - f, per image and append it to an error list
420     for idx in range(len(R_list)):
421         P = get_projection_matrix(K, R_list[idx], t_list[idx])
422         reprojected_pts = apply_homography(pts_3d, P)
423         # always the geometric error here
424         err = np.abs((pts_img_list[idx] - reprojected_pts)).ravel()
425         error = np.append(error, err, axis=0)
426
427     return error
428
429 def reprojection_error(pts_3d, pts_img, P):
430     # this is just another function to use to give us the same thing but
431     # this is just for one image
432     reprojected_pts = apply_homography(pts_3d, P)
433     return np.abs((pts_img - reprojected_pts)).ravel()
434
435 def lm_optim_P(pts_3d, pts_img_list, K, R_list, t_list):
436     # this is our optimization function, we call scipy least squares here
437     # after getting all the parameters
438     params = get_params_from_matrices(K, R_list, t_list)
439     op_params = least_squares(error_f, params, method='lm', args=(pts_3d,
440     pts_img_list), verbose=False).x
441     # we reconstruct them and this is what we return
442     K_op, R_op_list, t_op_list = get_matrices_from_params(op_params)
443
444     return K_op, R_op_list, t_op_list
445
446 def get_camera_poses(R, t):
447     # this is how we get our camera poses,
448     # there is one modification we did, we are no longer adding the C into
449     # the directions,
450     # as the function we call to draw the vectors in 3d gets the direction
451     # not the endpoint
452     C = - R.T @ t
453
454     X_xcam = np.array([1,0,0])

```

```
450 X_ycam = np.array([0,1,0])
451 X_zcam = np.array([0,0,1])
452
453 X_x = R.T @ X_xcam
454 X_y = R.T @ X_ycam
455 X_z = R.T @ X_zcam
456
457
458 return C, X_x, X_y, X_z
459
460 def plot_corners(image, corners_, name):
461     # this is just to plot corners and the label
462     image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
463     img = image.copy()
464     cont = 0
465     ms = 2
466     plt.figure()
467     plt.imshow(img);
468     for corner in corners_:
469         x = corner[0]
470         y = corner[1]
471         plt.plot(x,y, 'g.', markersize=ms)
472         plt.text(x, y, str(cont), fontsize=8, color='g')
473         cont += 1
474     plt.axis('off');
475     plt.savefig(name[:-4] + ".png",bbox_inches='tight',pad_inches=None);
476     plt.close()
477     plt.cla()
478     plt.clf()
479
480 def plot_corners_improvement(image, corners_, corners_improved, name):
481     # we use this to plot the detected corners and the reprojected corners
482     # projected corners always in red and detected corners in green
483     image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
484     img = image.copy()
485     cont = 0
486     ms = 2
487     plt.figure()
488     plt.imshow(img);
489     for corner in corners_:
490         x = corner[0]
491         y = corner[1]
492         plt.plot(x,y, 'r.', markersize=ms)
493         #plt.text(x, y, str(cont), fontsize=8, color='g')
494         #cont += 1
495     for corner in corners_improved:
496         x = corner[0]
497         y = corner[1]
498         plt.plot(x,y, 'g.', markersize=ms)
499         plt.text(x, y, str(cont), fontsize=8, color='g')
500         cont += 1
501     plt.axis('off');
502     plt.savefig(name[:-4] + ".png",bbox_inches='tight',pad_inches=None);
503     plt.close()
```

```

504     plt.cla()
505     plt.clf()
506
507 def reproject_to_fixed_image(fixed_img, fixed_image_data, other_image_list
508 ):
509     #this is the function we use to call the reprojection into the fixed
510     image,
511     # we are only interested in the first, second and last column of the
512     projection matrix,
513     # so we take just that from the projection matrix and use the
514     resulting 3 by 3 matrix to project the points
515     P_fixed = fixed_image_data["projection_matrix"]
516     H_fixed = P_fixed[:,[0,1,3]]
517     corners_fixed = fixed_image_data["corners"]
518     P_fixed_op = fixed_image_data["projection_matrix_new"]
519     H_fixed_op = P_fixed_op[:,[0,1,3]]
520     save_path = "improvement"
521     for image_data in other_image_list:
522         print("-"*10)
523         print(image_data["filename"])
524         corners = image_data["corners"]
525         P_old = image_data["projection_matrix"]
526         H_old = P_old[:,[0,1,3]]
527
528         projection_3d_old = apply_homography(corners, np.linalg.inv(H_old))
529     )
530     reprojected_fixedim_corners_old = apply_homography(
531     projection_3d_old, H_fixed)
532
533     plot_corners_improvement(fixed_img,
534     reprojected_fixedim_corners_old, corners_fixed, os.path.join(save_path,
535     image_data["filename"][:-4]+"_beforeLM.jpg"))
536
537     P_op = image_data["projection_matrix_new"]
538     H_op = P_op[:,[0,1,3]]
539
540     projection_3d_op = apply_homography(corners, np.linalg.inv(H_op))
541     reprojected_fixedim_corners_new = apply_homography(
542     projection_3d_op, H_fixed_op)
543     plot_corners_improvement(fixed_img,
544     reprojected_fixedim_corners_new, corners_fixed, os.path.join(save_path,
545     image_data["filename"][:-4]+"_afterLM.jpg"))
546     err_old = np.linalg.norm(reprojected_fixedim_corners_old -
547     fixed_image_data["corners"],axis=1)
548
549     err_new = np.linalg.norm(reprojected_fixedim_corners_new -
550     fixed_image_data["corners"],axis=1)
551
552     print("Mean: ", err_old.mean())
553     print("Mean after LM: ", err_new.mean())
554     print("-"*4)
555     print("Std: ", err_old.std())
556     print("Std after LM: ", err_new.std())

```



```

545 def plot_camera_poses(data_dict_list, pattern_shape=(8,10), pattern_size=1.,
546     name=None):
547     # this is how we plot the camera poses, it takes all the data at once
548     # and loops through it
549     plt.close()
550     plt.cla()
551     plt.clf()
552     fig = plt.figure(figsize=(11, 11))
553     ax = fig.add_subplot(projection='3d')
554     # this is for the calibration pattern
555     for j in range(pattern_shape[1]-1):
556         for i in range(pattern_shape[0]-1):
557             # need to reflect the fact that it "skips" a line
558             if j % 2 == 0:
559                 # alternate black and white
560                 color = 'black' if (i + j) % 2 == 0 else 'white'
561                 x = i * pattern_size
562                 y = j * pattern_size
563                 rect = Rectangle((x, y), pattern_size, pattern_size, color
564                     =color, alpha=1, edgecolor=None)
565                 ax.add_patch(rect)
566                 # Set the 3D position of each rectangle at Z=0
567                 pathpatch_2d_to_3d(rect, z=0, zdir="z")
568             else:
569                 # in this case, it is all white squares
570                 color = 'white'
571                 x = i * pattern_size
572                 y = j * pattern_size
573                 rect = Rectangle((x, y), pattern_size, pattern_size, color
574                     =color, alpha=1, edgecolor=None)
575                 ax.add_patch(rect)
576                 # Set the 3D position of each rectangle at Z=0
577                 pathpatch_2d_to_3d(rect, z=0, zdir="z")
578     # now we draw our poses
579     for img_data_dict in data_dict_list:
580         # loop through the data, choose a random color, the way matplotlib
581         # does it is with 0 to 1
582         color_random = np.random.random(3)
583         center = img_data_dict["camera_pose"]["center"]
584         c_x, c_y, c_z = center
585         x_dir = img_data_dict["camera_pose"]["x_vector"]
586         x_x, x_y, x_z = x_dir
587         y_dir = img_data_dict["camera_pose"]["y_vector"]
588         y_x, y_y, y_z = y_dir
589         z_dir = img_data_dict["camera_pose"]["z_vector"]
590         z_x, z_y, z_z = z_dir
591
592         # this is for our direction vectors
593         ax.quiver(c_x, c_y, c_z, x_x, x_y, x_z, color='r', linewidth=1)
594         ax.quiver(c_x, c_y, c_z, y_x, y_y, y_z, color='g', linewidth=1)
595         ax.quiver(c_x, c_y, c_z, z_x, z_y, z_z, color='b', linewidth=1)
596
597     # now the camera plane, in here we notice that the corners are
598     # made up from addition and substractions of the camera center and the

```

```
directions,
    # so we make use to that
    points = [center + x_dir + y_dir,
               center + x_dir - y_dir,
               center - x_dir - y_dir,
               center - x_dir + y_dir]

    poly = Poly3DCollection([points], color=color_random, alpha=0.4,
                             edgecolor=None)
    ax.add_collection3d(poly)

    #ax.view_init(45,-90)
    # we set reasonable limits and plot it
    ax.set_xlim([-15, 15])
    ax.set_ylim([-15, 15])
    ax.set_xlabel("x")
    ax.set_ylabel("y")
    ax.set_zlabel("z")
    if name:
        plt.savefig("camera_poses_.pdf",bbox_inches='tight',pad_inches=
None);
    else:
        plt.savefig("camera_poses.pdf",bbox_inches='tight',pad_inches=None
);
    plt.close()
    plt.cla()
    plt.clf()

img_path = "HW8-Files/Dataset1"
save_path = "imgs"
save_canny = "canny"
save_lines = "lines"

# we get all the files in the dataset to loop through
file_list = [x for x in os.listdir(img_path) if x.endswith(".jpg")]
imgs_data = []
H_list = []
corners_list = []
Vs = []

# again, 1 for length as we are using inches
RL_coords = get_coordinates(shape=(8,10),length=1)
# this is now our 3d coordinate list
RL_3d = add_number(RL_coords,0)

for file in file_list:
    img_data = {}
    img_data["filename"] = file
    img = cv2.imread(os.path.join(img_path,file))
    sname = os.path.join(save_path, file)
    canny_name = os.path.join(save_canny, file)
    line_name = os.path.join(save_lines, file)
    hough_name = os.path.join('houghlines', file)
    corners = get_corners(img, tol=15, save_canny=canny_name, save_lines=
line_name, save_hough=hough_name)
    corners_list.append(corners)
```

```

642     img_data["corners"] = np.array(corners)
643
644
645     plot_corners(img, corners, sname);
646
647     H = get_H_matrix(RL_coords, np.array(corners))
648     img_data["homography"] = H
649     #print(H)
650     if len(corners) == 80:
651         # need to make sure that the number of corners is 80
652         H_list.append(H)
653
654         V_h = get_V(H)
655         Vs.append(V_h[0])
656         Vs.append(V_h[1])
657
658         imgs_data.append(img_data)
659     else:
660         print("Incorrect number of lines for ", file)
661
662     Vs = np.array(Vs)
663     omega = get_omega(Vs)
664     K = get_K_matrix(omega)
665
666     R_list = []
667     t_list = []
668
669     for idx in range(len(imgs_data)):
670         R, t = get_extrinsics(K, imgs_data[idx]["homography"])
671         imgs_data[idx]["rotation_matrix"] = R
672         imgs_data[idx]["translation_vector"] = t
673         imgs_data[idx]["projection_matrix"] = get_projection_matrix(K, R, t)
674         R_list.append(R)
675         t_list.append(t)
676     # now the optimization
677     K_op, R_op_list, t_op_list = lm_optim_P(RL_3d, corners_list, K, R_list,
        t_list)
678
679     for idx in range(len(imgs_data)):
680         # loop through the data and populate what is missing
681         imgs_data[idx]["rotation_matrix"] = R_list[idx]
682         imgs_data[idx]["translation_vector"] = t_list[idx]
683         imgs_data[idx]["projection_matrix"] = get_projection_matrix(K, R_list[
            idx], t_list[idx])
684
685         imgs_data[idx]["K_op"] = K_op
686         imgs_data[idx]["R_op"] = R_op_list[idx]
687         imgs_data[idx]["t_op"] = t_op_list[idx]
688
689         P_new = get_projection_matrix(K_op, R_op_list[idx], t_op_list[idx])
690         P_old = get_projection_matrix(K, R_list[idx], t_list[idx])
691         imgs_data[idx]["projection_matrix_new"] = P_new
692

```

```

693     imgs_data[idx]["old_reprojection_error"] = reprojection_error(RL_3d,
694     imgs_data[idx]["corners"], P_old)
695     imgs_data[idx]["new_reprojection_error"] = reprojection_error(RL_3d,
696     imgs_data[idx]["corners"], P_new)
697     imgs_data[idx]["err_mean_old"] = imgs_data[idx]["
698     old_reprojection_error"].mean()
699     imgs_data[idx]["err_mean_new"] = imgs_data[idx]["
700     new_reprojection_error"].mean()
701     imgs_data[idx]["err_std_old"] = imgs_data[idx]["old_reprojection_error
702     "].std()
703     imgs_data[idx]["err_std_new"] = imgs_data[idx]["new_reprojection_error
704     "].std()
705
706     C, X_x, X_y, X_z = get_camera_poses(R_op_list[idx], t_op_list[idx])
707     camera_pose = {}
708     camera_pose["center"] = C
709     camera_pose["x_vector"] = X_x
710     camera_pose["y_vector"] = X_y
711     camera_pose["z_vector"] = X_z
712     imgs_data[idx]["camera_pose"] = camera_pose
713
714     print("K before LM: ", K)
715     print("K after LM: ", K_op)
716
717     imgs_to_show = ["Pic_19.jpg", "Pic_31.jpg"]
718     fixed_img = "Pic_11.jpg"
719     fixed_image = cv2.imread(os.path.join(img_path, fixed_img))
720     imgs_show_data = []
721     for img_data in imgs_data:
722         if img_data["filename"] == fixed_img:
723             fixed_img_data = img_data
724         if img_data["filename"] in imgs_to_show:
725             imgs_show_data.append(img_data)
726
727     reproject_to_fixed_image(fixed_image, fixed_img_data, imgs_show_data)
728
729     for img_show_data in imgs_show_data:
730         print("-"*10)
731         print(img_show_data["filename"])
732         print("Before LM, rotation matrix: ", img_show_data["rotation_matrix"])
733         print("Before LM, translation vector: ", img_show_data["
734         translation_vector"])
735         print("After LM, rotation matrix: ", img_show_data["R_op"])
736         print("After LM, translation vector: ", img_show_data["t_op"])
737
738         img = cv2.imread(os.path.join(img_path, img_show_data['filename']))
739         new_points = apply_homography(RL_3d, img_show_data["projection_matrix"]
740         ])
741         new_points_improved = apply_homography(RL_3d, img_show_data["
742         projection_matrix_new"])
743         plot_corners_improvement(img, new_points, img_show_data["corners"], os
744         .path.join("improvement", img_show_data["filename"][: -4] + "
745         _corners_beforeLM.jpg"))

```

```
735     plot_corners_improvement(img, new_points_improved, img_show_data["  
corners"], os.path.join("improvement",img_show_data["filename"][:-4]+"  
_corners_afterLM.jpg"))  
736     plot_corners_improvement(img, new_points, new_points_improved, os.path  
.join("improvement",img_show_data["filename"][:-4]+"_corners_comparison  
.jpg"))  
737  
738 plot_camera_poses(imgs_data)  
739  
740 # for the second dataset we follow exactly the same steps  
741 img_path_ = "dataset2/small"  
742 save_path = "imgs"  
743 save_canny = "canny"  
744 save_lines = "lines"  
745 file_list_ = [x for x in os.listdir(img_path_) if x.endswith(".jpeg")]  
746 imgs_data_ = []  
747 H_list_ = []  
748 corners_list_ = []  
749 Vs_ = []  
750 RL_coords = get_coordinates(shape=(8,10),length=1)  
751 RL_3d = add_number(RL_coords,0)  
752  
753 for file in file_list_:  
754     img_data_ = {}  
755     img_data_["filename"] = file  
756     img = cv2.imread(os.path.join(img_path_,file))  
757     sname = os.path.join(save_path, file)  
758     canny_name = os.path.join(save_canny, file)  
759     line_name = os.path.join(save_lines, file)  
760     hough_name = os.path.join('houghlines', file)  
761     corners = get_corners(img, tol=20, save_canny=canny_name, save_lines=  
line_name, save_hough=hough_name, canny_p1=370, canny_p2=300, hough_p  
=50)  
762     corners_list_.append(corners)  
763     img_data_["corners"] = np.array(corners)  
764  
765     if len(corners) == 80:  
766  
767         plot_corners(img, corners, sname);  
768  
769         H = get_H_matrix(RL_coords, np.array(corners))  
770         img_data_["homography"] = H  
771         #print(H)  
772         H_list_.append(H)  
773  
774         V_h = get_V(H)  
775         Vs_.append(V_h[0])  
776         Vs_.append(V_h[1])  
777  
778         imgs_data_.append(img_data_)  
779     else:  
780         print(file)  
781         print(len(corners))  
782
```

```

783 Vs_ = np.array(Vs_)
784 omega = get_omega(Vs_)
785 K_ = get_K_matrix(omega)
786
787 R_list_ = []
788 t_list_ = []
789
790 for idx in range(len(imgs_data_)):
791     R, t = get_extrinsics(K_, imgs_data_[idx]["homography"])
792     imgs_data_[idx]["rotation_matrix"] = R
793     imgs_data_[idx]["translation_vector"] = t
794     imgs_data_[idx]["projection_matrix"] = get_projection_matrix(K_, R, t)
795     R_list_.append(R)
796     t_list_.append(t)
797
798 K_op_, R_op_list_, t_op_list_ = lm_optim_P(RL_3d, corners_list_, K_,
799     R_list_, t_list_)
800
801 for idx in range(len(imgs_data_)):
802     imgs_data_[idx]["rotation_matrix"] = R_list_[idx]
803     imgs_data_[idx]["translation_vector"] = t_list_[idx]
804     imgs_data_[idx]["projection_matrix"] = get_projection_matrix(K_,
805     R_list_[idx], t_list_[idx])
806
807     imgs_data_[idx]["K_op"] = K_op_
808     imgs_data_[idx]["R_op"] = R_op_list_[idx]
809     imgs_data_[idx]["t_op"] = t_op_list_[idx]
810
811     P_new = get_projection_matrix(K_op_, R_op_list_[idx], t_op_list_[idx])
812     P_old = get_projection_matrix(K_, R_list_[idx], t_list_[idx])
813     imgs_data_[idx]["projection_matrix_new"] = P_new
814
815     imgs_data_[idx]["old_reprojection_error"] = reprojection_error(RL_3d,
816     imgs_data_[idx]["corners"], P_old)
817     imgs_data_[idx]["new_reprojection_error"] = reprojection_error(RL_3d,
818     imgs_data_[idx]["corners"], P_new)
819     imgs_data_[idx]["err_mean_old"] = imgs_data_[idx]["old_reprojection_error"].mean()
820     imgs_data_[idx]["err_mean_new"] = imgs_data_[idx]["new_reprojection_error"].mean()
821     imgs_data_[idx]["err_std_old"] = imgs_data_[idx]["old_reprojection_error"].std()
822     imgs_data_[idx]["err_std_new"] = imgs_data_[idx]["new_reprojection_error"].std()
823
824     C, X_x, X_y, X_z = get_camera_poses(R_op_list_[idx], t_op_list_[idx])
825     camera_pose = {}
826     camera_pose["center"] = C
827     camera_pose["x_vector"] = X_x
828     camera_pose["y_vector"] = X_y
829     camera_pose["z_vector"] = X_z
830     imgs_data_[idx]["camera_pose"] = camera_pose
831
832 print("K before LM: ", K_)

```

```

829 print("K after LM: ", K_op_)
830
831 imgs_to_show = ["pose9.jpeg", "pose33.jpeg"]
832 fixed_img = "fixed_img.jpeg"
833 fixed_image = cv2.imread(os.path.join(img_path_, fixed_img))
834 imgs_show_data = []
835 for img_data in imgs_data_:
836     if img_data["filename"] == fixed_img:
837         fixed_img_data = img_data
838     if img_data["filename"] in imgs_to_show:
839         imgs_show_data.append(img_data)
840
841 reproject_to_fixed_image(fixed_image, fixed_img_data, imgs_show_data)
842
843 print(fixed_img_data["rotation_matrix"])
844 print(fixed_img_data["translation_vector"])
845 print(fixed_img_data["R_op"])
846 print(fixed_img_data["t_op"])
847
848 for img_show_data in imgs_show_data:
849     print("-"*10)
850     print(img_show_data["filename"])
851     print("Before LM, rotation matrix: ",img_show_data["rotation_matrix"])
852     print("Before LM, translation vector: ",img_show_data["translation_vector"])
853     print("After LM, rotation matrix: ",img_show_data["R_op"])
854     print("After LM, translation vector: ",img_show_data["t_op"])
855
856     img = cv2.imread(os.path.join(img_path_, img_show_data['filename']))
857     new_points = apply_homography(RL_3d, img_show_data["projection_matrix"]
858 ]
859     new_points_improved = apply_homography(RL_3d, img_show_data["projection_matrix_new"])
860     plot_corners_improvement(img, new_points, img_show_data["corners"], os
861 .path.join("improvement",img_show_data["filename"][: -4]+ "_corners_beforeLM.jpg"))
862     plot_corners_improvement(img, new_points_improved, img_show_data["corners"], os.path.join("improvement",img_show_data["filename"][: -4]+ "_corners_afterLM.jpg"))
863     plot_corners_improvement(img, new_points, new_points_improved, os.path
864 .join("improvement",img_show_data["filename"][: -4]+ "_corners_comparison.jpg"))
865
866 plot_camera_poses(imgs_data_, name=True)

```

Listing 1: Source code