

# Homework 10

## Theory Question

1. Overfitting to the training data means that our model learns only the training data on a superficial level, and not the representations underneath that best describe our dataset. For example, in linear regression, if our classifier overfits, that means that it has not learned what is underneath but has learned some curve with passes through all the data points.
2. The problem is that if we try to sample from the latent space, this procedure is not differentiable. So to make it differentiable there is a need to change this sampling procedure. What is done is to not sample from the distribution  $z \sim \mathcal{N}(\mu, \sigma^2)$  but we sample from a zero-mean unit variance normal  $\epsilon \sim \mathcal{N}(0, 1)$ . So now we add our mean  $\mu$  and multiply it by our variance. So now,  $z = \mu + \epsilon\sigma$ . Now this operation is differentiable, so essentially we change the way we sample from the normal, through making use of the zero mean unit variance normal distribution.

## Task 1

For this and Task 2 we will be using the FacePix dataset provided with the homework files. At the very start, we associate each image to a label represented by the first part of the image file's name. In this case we have 30 different classes (30 faces of 30 different people).

## PCA

We begin by vectorizing each image, turning it from, in this case of size  $128 \times 128$ , to a vector of length 16384. Then we can compute a mean image given by:

$$\mathbf{m} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i \quad (1)$$

Where  $N$  is the number of images and  $\mathbf{x}_i$  is the vectorized  $i^{\text{th}}$  image. With this we can now compute the covariance matrix. The covariance matrix is given by:

$$\mathbf{C} = \frac{1}{N} \sum_{i=1}^N \{(\mathbf{x}_i - \mathbf{m})(\mathbf{x}_i - \mathbf{m})^T\} \quad (2)$$

After subtracting the mean from the vectorized images, we normalize them each, ensuring that the norm of each image/row is equal to 1. The next step is to get the SVD decomposition of this covariance matrix, however, we run into issues since this covariance matrix is of size  $16384 \times 16384$ , making it computational resource intensive to be able to do this decomposition. We make use of the computational trick to alleviate this. Our new covariance matrix is now:

$$\mathbf{C} = \frac{1}{N} \sum_{i=1}^N \{(\mathbf{x}_i - \mathbf{m})^T(\mathbf{x}_i - \mathbf{m})\} \quad (3)$$

Which results in a  $N \times N$  size covariance matrix. We perform SVD decomposition on this smaller size covariance matrix. Now, we are interested in the eigenvectors of the bigger matrix, given by  $\mathbf{w}$ . We know that our smaller matrix eigenvectors are related to these by:

$$\mathbf{w} = \mathbf{X}\mathbf{v} \quad (4)$$

Where  $\mathbf{X}$  is the entire stack of images and  $\mathbf{v}$  is the eigenvectors from the smaller covariance matrix. Now, all that is left is for us to normalize  $\mathbf{W}$ , and to pick the first  $p$  eigenvectors. Now that we have the eigenvectors, we can compute the projection of our images into the eigenvectors. This projection is now our training "features".

**Nearest Neighbor** To be able to classify our faces, we make use of nearest neighbor algorithm. This consists of getting the distance from a data point to all the points of the training set, we then take the closest training set data point and assign our test dataset point the same label as this closest training set data point.

For the testing set, we follow the same process. First, we retain the mean image, that is calculated from the training set, and we subtract it from the test data as well. Then we get the projection of all our testing set data points to the space generated by the eigenvectors  $\mathbf{w}$ . We use nearest neighbor to classify these testing data points.

**UMAP Embeddings** The training and testing features undergo dimensionality reduction with UMAP library, resulting in only 2 dimensions. For the training set, the classes are set with the ground truth labels and for the test set the classes are indicated by the predicted classes. We plot these embeddings in a 2D graph.

## LDA

We begin similarly to PCA. The main difference is that we now need to divide the images into their respective classes. We now calculate the per class mean given by:

$$\mathbf{m}_i = \frac{1}{|\mathcal{C}_i|} \sum_{n=1}^{|\mathcal{C}_i|} \mathbf{x}_n \quad (5)$$

Where  $\mathcal{C}_i$  is the  $i^{\text{th}}$  class. We then make use of all the class means and the global mean to get the between class scatter and the within class scatter, given by:

$$\mathbf{S}_B = \frac{1}{|\mathcal{C}|} \sum_{i=1}^{|\mathcal{C}|} \{(\mathbf{m}_i - \mathbf{m})^T(\mathbf{m}_i - \mathbf{m})\} \quad (6)$$

$$\mathbf{S}_W = \frac{1}{|\mathcal{C}|} \sum_{i=1}^{|\mathcal{C}|} \frac{1}{|\mathcal{C}_i|} \sum_{n=1}^{|\mathcal{C}_i|} \{(\mathbf{x}_n - \mathbf{m}_i)^T(\mathbf{x}_n - \mathbf{m}_i)\} \quad (7)$$

Where  $\mathbf{S}_B$  is the between class scatter matrix and  $\mathbf{S}_W$  is the within class scatter matrix. We apply the Yu-Yang algorithm, which consists of starting with a similar step to the

computational trick for calculating the between class scatter matrix. We obtain  $\mathbf{S}_B$  through:

$$\mathbf{S}_B = \frac{1}{|\mathcal{C}|} \sum_{i=1}^{|\mathcal{C}|} \{(\mathbf{m}_i - \mathbf{m})(\mathbf{m}_i - \mathbf{m})^T\} \quad (8)$$

Which gives us the same number of eigenvectors as images in our training set after following the same steps to get the actual eigenvectors similar to PCA. Then we calculate a matrix  $\mathbf{Z}$ , given by:

$$\mathbf{Z} = \mathbf{w}\mathbf{D}^{-1/2} \quad (9)$$

Where  $\mathbf{D}$  is the diagonal matrix built from the eigenvalues and  $\mathbf{w}$  is the eigenvectors obtained previously. One more thing to consider in this step is that we are discarding the smallest 5 eigenvalues and their corresponding eigenvectors. To get the eigenvectors of  $\mathbf{S}_w$  we make use of  $\mathbf{Z}$  and the computational trick.

$$\mathbf{Z}_X = \mathbf{Z}^T \mathbf{X} \quad (10)$$

Now we can get  $\mathbf{S}_w$  from this:

$$\mathbf{S}_w = \mathbf{Z}_X \mathbf{Z}_X^T \quad (11)$$

Now we get the eigenvectors of  $\mathbf{S}_w$  and only retain  $p$  of them, similar to PCA. Finally, the classification step is done with the same steps as for PCA, the data points are all projected to the space of the eigenvectors, and then nearest neighbor is used to classify the test set data points. UMAP embeddings are also obtained for this.

## Task 2: Autoencoder

We make the equivalence of the autoencoder weights at epoch  $e$  with the  $p$  chosen eigenvectors from PCA and LDA so we can graph both of them in one single graph. With the given autoencoder weights, we get the latent space variables for each of images in the training and testing set, and it is this latent space variables that we call our features. We then use the previously mentioned nearest neighbor to classify get the closest train feature and classify it with that label. We repeat this process for all 3 weights given: 3, 8 and 16 epochs. We also do dimensionality reduction to get the UMAP embeddings and plot them.

## Task 3: Adaboost Cascade

### Feature extraction

To begin with, we extract horizontal and vertical features with a 1D Haar filter. These filters look like:  $(-1 \ -1 \ -1 \ 1 \ 1 \ 1)$  for horizontal and  $(-1 \ -1 \ -1 \ 1 \ 1 \ 1)^T$  for vertical features. We convolve the image with several of these filters, with sizes ranging from 2 to the width/height of the image. We append every result of the convolution to a vector, we treat these as our features. We repeat this process for every image in both the training and testing set.

## Weak Classifiers

Now that we have all the features for our training set and a given set of weights, we proceed to loop through all the features. We now sort them, and find the error:

$$\epsilon_1 = S^+ + T^- - S^- \quad (12)$$

However, we need to consider the inverse polarity for our thresholding, so we consider another error:

$$\epsilon_2 = S^- + T^+ - S^+ \quad (13)$$

In both,  $S$  is the cumulative sum of the weights for the positive or negative samples, and  $T$  is the total sum of the positive or negative weights. We now find the minimum of each error and find whether this minimum belongs to  $\epsilon_1$  or  $\epsilon_2$ , as this will determine the polarity for our thresholding. We find the feature where the minimum error is in and store it. We compare this minimum error value to a minimum value found across all features to only keep the best weak classifier. Once this best weak classifier is found, we save the polarity, threshold value, which feature is involved, and the predictions on the training set.

## Cascade/Strong classifier

We initialize the weights, making sure that positive and negative labels have collective weights of  $1/2$ . All the positive labels will then have the same weight, so they are normalized depending on how many there are, same for negative labels. Now we take the entire training feature set and obtain a weak classifier. We then take the predictions from the weak classifier and the error from this weak classifier to update our weights for the next weak classifier.

$$\begin{aligned} \beta_t &= \frac{\epsilon_t}{1 - \epsilon_t} \\ \alpha_t &= \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right) \end{aligned} \quad (14)$$

We use this to update the weights for our next weak classifier:

$$\begin{aligned} D_{t+1}(x_i) &= \frac{D_t(x_i) \{\beta_t(h(x_i) - y_i)\}}{Z_t} \\ Z_t &= \sum_{i=1}^m D_t(x_i) \{\beta_t(h(x_i) - y_i)\} \end{aligned} \quad (15)$$

As for the trust factor,  $\alpha_t$ , we will use it for the final predictions. We repeat this process for as many classifiers we want our cascade to have. Our final classifier for the cascade will be given by:

$$H(x_i) = \sum_{t=1}^N \alpha_t h_t(x_i) \geq \sum_{t=1}^N \frac{1}{2} \alpha_t \quad (16)$$

## Training

For training, we get our first strong classifier, then depending on the final predictions of this classifier, we revise the training set. We keep all the positive samples, but for the negative samples, we only keep those that have been misclassified. We revise both the features and the labels for the next cascade. We repeat this for as many cascades as desired.




## Inference


For inference, we take the features from the testing set and put them through the first cascade. We obtain the final predictions, and in a similar manner to when we were training, we need to revise the testing set. Our revision consists of keeping only the positive predicted features. After this, we put it through the next cascade and so on until the final cascade.

## Results


### Task 1 & 2: PCA, LDA and Autoencoder

As we can see in Fig. 1, LDA outperforms PCA in every single value of  $p$  tested. It also achieved 100% classification accuracy, while PCA did not. We also see that PCA tends to not always have an increase in accuracy as the value of  $p$  increases, as we can see for  $p = 7$ ,  $p = 13$ . In contrast, LDA does increase in as the value of  $p$  increases. Now, as for the autoencoder, it starts off about as good as LDA, and better than PCA, however it quickly loses ground to LDA, as at  $p = 8$  it is worse than LDA, but better than PCA. It also reaches 100% accuracy, which PCA was not able to do. As for the UMAP embeddings, we can see how as we increase  $p$  these tend to be drawn together, in low  $p$  values these are very spread out and as a result our nearest neighbor misclassifies them. This is common to all 3: PCA, LDA and autoencoder. As we can see, the best performing classifier is LDA, and for its  $p = 15$  embeddings we can see that these are closely bunched together, compared to PCA and autoencoder.


Figure 1: Accuracy as a function of  $p$ .(a) PCA UMAP embeddings for  $p = 5$ , training set.(b) PCA UMAP embeddings for  $p = 5$ , testing set.




(c) PCA UMAP embeddings for  $p = 15$ , training set.




(d) PCA UMAP embeddings for  $p = 15$ , testing set.




(e) LDA UMAP embeddings for  $p = 5$ , training set.




(f) LDA UMAP embeddings for  $p = 5$ , testing set.




(g) LDA UMAP embeddings for  $p = 15$ , training set.



(h) LDA UMAP embeddings for  $p = 15$ , testing set.




(i) Autoencoder UMAP embeddings for  $p = 3$ , training set.




(j) Autoencoder UMAP embeddings for  $p = 3$ , testing set.





(k) Autoencoder UMAP embeddings for  $p = 16$ , training set.



(l) Autoencoder UMAP embeddings for  $p = 16$ , testing set.

### Task 3: Adaboost Cascade




Figure 2: FPR and FNR during training, as a function of the stages.




Figure 3: FPR and FNR during training, as a function of the stages.

## Source code

```

1  import cv2
2  import os
3  import torch
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import umap.umap_ as umap
7  from autoencoder import *
8
9
10 train_path = "FaceRecognition/train"
11 train_list = [x for x in os.listdir(train_path) if x.endswith(".png")]
12 train_list = sorted(train_list)
13
14 test_path = "FaceRecognition/test"
15 test_list = [x for x in os.listdir(test_path) if x.endswith(".png")]
16 test_list = sorted(test_list)
17
18 def load_data_autoencoder(TRAIN_DATA_PATH, EVAL_DATA_PATH, p):
19     # this is just to load the autoencoder, and then use it to obtain the
20     # latent variables for each image,
21     # these now become our training features/testing features
22     model = Autoencoder(p)

```

```

22     LOAD_PATH = f'weights/model_{p}.pt'
23     trainloader = DataLoader(
24         dataset=DataBuilder(TRAIN_DATA_PATH),
25         batch_size=1,
26     )
27     model.load_state_dict(torch.load(LOAD_PATH))
28     model.eval()
29
30     X_train, y_train = [], []
31     for batch_idx, data in enumerate(trainloader):
32         mu, logvar = model.encode(data['x'])
33         z = mu.detach().cpu().numpy().flatten()
34         X_train.append(z)
35         y_train.append(data['y'].item())
36     X_train = np.stack(X_train)
37     y_train = np.array(y_train)
38
39     testloader = DataLoader(
40         dataset=DataBuilder(EVAL_DATA_PATH),
41         batch_size=1,
42     )
43     X_test, y_test = [], []
44     for batch_idx, data in enumerate(testloader):
45         mu, logvar = model.encode(data['x'])
46         z = mu.detach().cpu().numpy().flatten()
47         X_test.append(z)
48         y_test.append(data['y'].item())
49     X_test = np.stack(X_test)
50     y_test = np.array(y_test)
51
52     return X_train, y_train, X_test, y_test
53
54 def get_img_data(file_list, img_path):
55     # this is just to obtain the flattened images and their corresponding
56     # label,
57     # we obtain the label from the name of the image file
58     mat = []
59     labels = []
60     for file in file_list:
61         img = cv2.imread(os.path.join(img_path, file), cv2.IMREAD_GRAYSCALE)
62         labels.append(int(file[:2]))
63         mat.append(img.reshape(-1))
64     return np.array(mat), np.array(labels)
65
66 def normalize_matrix(img_data):
67     # we normalize the rows
68     norm_ = np.linalg.norm(img_data, axis=1)[:, None]
69     n_imgdata = img_data/norm_
70     return n_imgdata
71
72 def get_covariance(img_data):
73     # this is to get the covariance matrix
74     covariance = img_data.T @ img_data
75     covariance /= (len(covariance) - 1)

```

```

75     return covariance
76
77 def get_pca(img_data, p=10):
78     # subtract the global mean, then normalize the data
79     img_data = img_data - img_data.mean(axis=0)[None,...]
80     normed_data = normalize_matrix(img_data)
81     normed_dataT = normed_data.T
82     # get the covariance then get the svd decomposition, proceed with the
    computational trick and pick only the p first eigenvectors
83     cov = get_covariance(normed_dataT)
84     _,_, v = np.linalg.svd(cov)
85     # up next is the computational trick
86     # I transpose the matrix here to reuse our normalization function,
87     # also because I understand it better if it has shape (N,16000) over
    (16000,N)
88     # this is seen in the LDA function next up
89     W = (normed_dataT @ v).T
90     normed_W = normalize_matrix(W)
91     return normed_W[:p]
92
93 def get_lda(img_data, labels, p=10):
94     class_mean = {}
95     per_class = {}
96     # we get the class means and put them into that dictionary
97     global_mean = img_data.mean(axis=0)[None,...]
98     C_mean = []
99     for id_class in range(labels.min(), labels.max()+1):
100         # we append the mean image per class to a matrix that has all of
    them
101         cls_img_data = img_data[labels == id_class]
102         per_class[str(id_class)] = cls_img_data.shape[0]
103         class_mean[str(id_class)] = cls_img_data.mean(axis=0)[None,...]
104         C_mean.append(cls_img_data.mean(axis=0))
105     # we use the yu-yang algorithm and also the computational trick in
    here
106     C_mean = np.array(C_mean)
107     C_mean_m = C_mean - global_mean
108     img_data = img_data - global_mean
109     #print(C_mean_m.shape)
110     S_b_c = C_mean_m @ C_mean_m.T
111     #print(S_b_c.shape)
112     _, d, v = np.linalg.svd(S_b_c)
113     # get the transpose to reuse our function, and also we ignore the last
    5 eigenvalues and eigenvectors, ignoring the last 5 worked well
114     w = (C_mean_m.T @ v).T
115     normed_w = normalize_matrix(w)[:5]
116     diagd = np.diag(d[:5])
117     z = np.linalg.inv(diagd) @ normed_w
118     # again I keep the matrices of shape (N,16000) as it is easier for me
    to grasp
119     z_x = z @ img_data.T
120     new_s = z_x @ z_x.T
121     _, _, v2 = np.linalg.svd(new_s)
122     new_w = (z.T @ v2).T

```

```
123     # print(new_w.shape)
124     return new_w[:p]
125
126
127 def NN(xprojection_train, y_train, xprojection_test):
128     # implementation of nearest neighbor, we make use of broadcasting to
129     # keep it to only one line of code
130     distances = np.linalg.norm(xprojection_test[:,None,:] -
131     xprojection_train[None,...], axis=2)
132     # get the index where the minimum happens
133     idx = np.argmin(distances,axis=1)
134     # return the label corresponding to it
135     return y_train[idx]
136
137
138 def get_accuracy(y_test, y_pred):
139     # get the accuracy as described in the homework pdf
140     correct = (y_test == y_pred).sum()
141     return correct/y_test.shape[0]
142
143
144 def plot_umap_embeddings(x_train, y_train, x_test,y_pred, mode='pca', p
145 =10, num_classes=30):
146     # plot the umap embeddings and save the plot
147     umap_t = umap.UMAP(n_components=2, random_state=0)
148     x_train_emb = umap_t.fit_transform(x_train)
149     x_test_emb = umap_t.transform(x_test)
150
151     colors = np.random.random((num_classes,3))
152
153     # train embeddings
154     fig = plt.figure(figsize=(11,11))
155     for id in range(1,num_classes):
156         x_train_emb_c = x_train_emb[y_train == id]
157         plt.scatter(x_train_emb_c[:,0],x_train_emb_c[:,1], color=colors[id
158 ])
159     plt.savefig(f'imgs/umap_emb_{mode}_p{p}_train.png', bbox_inches='tight
160 ', pad_inches=0)
161
162     plt.close()
163     plt.clf()
164     # test embeddings
165     fig = plt.figure(figsize=(11,11))
166     for id in range(1,num_classes):
167         x_test_emb_c = x_test_emb[y_pred == id]
168         plt.scatter(x_test_emb_c[:,0],x_test_emb_c[:,1], color=colors[id])
169     plt.savefig(f'imgs/umap_emb_{mode}_p{p}_test.png', bbox_inches='tight
170 ', pad_inches=0)
171     plt.close()
172     plt.clf()
173
174 def classifier(x_train, y_train, x_test, y_test, mode='pca', p=10,
175 plot_embeddings=False):
```

```
169 # this is to build the classifier with the training and testing
    features as well as plot the umap embeddings
170 g_mean = x_train.mean(axis=0)[None,...]
171 # we set 3 modes, LDA PCA and autoencoder, to not have to remake other
    functions
172 if mode!="autoencoder":
173     # get the projections to the eigenvectors
174     if mode=='pca':
175         eigs = get_pca(x_train, p=p)
176     elif mode == 'lda':
177         eigs = get_lda(x_train, y_train, p=p)
178     x_train = x_train - g_mean
179     x_test = x_test - g_mean
180
181     xp_train = (eigs @ x_train.T).T
182     xp_test = (eigs @ x_test.T).T
183 else:
184     # in the case of the autoencoder we don't need to do anything, as
    we already got our training features from the autoencoder itself
185     xp_train = x_train
186     xp_test = x_test
187     # we use nearest neighbor
188     y_pred = NN(xp_train, y_train, xp_test)
189     # get the accuracy
190     acc = get_accuracy(y_test, y_pred)
191
192     if plot_embeddings:
193         plot_umap_embeddings(xp_train, y_train, xp_test, y_pred, mode=mode,
            p=p)
194     return acc
195
196 x_train, y_train = get_img_data(train_list, train_path)
197 x_test, y_test = get_img_data(test_list, test_path)
198 # values of p we will plot
199 ps = np.arange(1,21,1)
200 # save the accuracy of both
201 pca_acc = []
202 lda_acc = []
203
204 for p in ps:
205     plot_emb = False
206     if p % 5 == 0:
207         plot_emb = True
208
209     # pca
210     pca_a = classifier(x_train, y_train, x_test, y_test, mode='pca', p=p,
        plot_embeddings=plot_emb)
211     pca_acc.append(pca_a)
212
213     # lda
214     lda_a = classifier(x_train, y_train, x_test, y_test, mode='lda', p=p,
        plot_embeddings=plot_emb)
215     lda_acc.append(lda_a)
216
```

```
217 p_ep = np.array([3,8,16])
218 aenc_acc = []
219 for p in p_ep:
220     # this is for our autoencoder, we plot on every value of p since
    theres only 3
221     x_train, y_train, x_test, y_test = load_data_autoencoder(train_path,
    test_path, p)
222     a_acc = classifier(x_train, y_train, x_test, y_test, mode='autoencoder
    ', p=p, plot_embeddings=True)
223     aenc_acc.append(a_acc)
224
225
226 # plot accuracies as function of p
227 fig = plt.figure(figsize=(11,11))
228 plt.plot(ps,pca_acc, marker="o",label="PCA")
229 plt.plot(ps,lda_acc, marker="v",label="LDA")
230 plt.plot(p_ep,aenc_acc, marker="s",label="Autoencoder")
231
232 plt.xlim(0,21,1);
233 plt.ylim(0,1.05)
234 plt.xticks(ps)
235 plt.legend()
236 plt.savefig(f'imgs/accs_p.png', bbox_inches='tight', pad_inches=0)
237 plt.close()
238
239 # some functions here have been based on the 2022's best solutions
240
241 def get_features(image):
242     # image is grayscale
243     h, w = image.shape
244     image = (image/255).astype(float)
245     sizes_vert = np.arange(2, h, 2)
246     sizes_hor = np.arange(2, w, 2)
247     features = []
248     # vertical features
249     for f_size in sizes_vert:
250         # pad in vertical direction only
251         image_p = cv2.copyMakeBorder(image, f_size//2, f_size//2,0,0, cv2.
    BORDER_CONSTANT, 0)
252         for jdx in range(f_size//2,image_p.shape[0]-f_size//2):
253             for idx in range(image_p.shape[1]):
254                 neg = image_p[jdx - f_size//2:jdx, idx].sum()
255                 pos = image_p[jdx:jdx + f_size//2 + 1, idx].sum()
256                 feature = pos - neg
257                 features.append(feature)
258     # horizontal features
259     for f_size in sizes_hor:
260         # pad in horizontal direction only
261         image_p = cv2.copyMakeBorder(image, 0, 0, f_size//2, f_size//2,
    cv2.BORDER_CONSTANT, 0)
262         for jdx in range(image.shape[0]):
263             for idx in range(f_size//2,image_p.shape[1]-f_size//2):
264                 neg = image_p[jdx, idx - f_size//2:idx].sum()
265                 pos = image_p[jdx, idx:idx + f_size//2 + 1].sum()
```

```
266         feature = pos - neg
267         features.append(feature)
268     # now we have all our features
269     features = np.array(features)
270
271     return features
272
273 def get_car_data(path, num=1):
274     # this is just to load our images and get the features
275     file_list = [x for x in os.listdir(path) if x.endswith("png")]
276     feature_list = []
277     for file in file_list:
278         img = cv2.imread(os.path.join(path, file), cv2.IMREAD_GRAYSCALE)
279         features = get_features(img)
280         feature_list.append(features)
281     feature_list = np.array(feature_list)
282     # we just get the labels like this, for negative we just change it to
283     0
284     class_label = np.ones(len(file_list)) * num
285     return feature_list, class_label
286
287 train_pos_path = "CarDetection/train/positive"
288 train_neg_path = "CarDetection/train/negative"
289 test_pos_path = "CarDetection/test/positive"
290 test_neg_path = "CarDetection/test/negative"
291
292 # we make sure to maintain the distinction between positive and negative
293 x_train_pos, y_train_pos = get_car_data(train_pos_path, 1)
294 x_train_neg, y_train_neg = get_car_data(train_neg_path, 0)
295
296 x_test_pos, y_test_pos = get_car_data(test_pos_path, 1)
297 x_test_neg, y_test_neg = get_car_data(test_neg_path, 0)
298
299
300 def weak_classifier(features, labels, weights):
301     # this is to get a weak classifier,
302     # to start we set a high error, since it ensures that it will always
303     change in the first iteration
304     cls_error = np.inf
305
306     for idxs in range(features.shape[1]):
307         # loop through features and order them,
308         # also order the labels and the weights accordingly
309         feature = features[:, idxs]
310         idxs_sort = np.argsort(feature)
311         feature_sort = feature[idxs_sort]
312         labels_sort = labels[idxs_sort, 0]
313         weights_sort = weights[idxs_sort, 0]
314         # this is to get the errors, we just set the negative weights to 0
315         for the cumulative positive sum and viceversa
316         weights_pos = weights_sort.astype(float)
317         weights_pos[labels_sort == 0] = 0
```



```

317     weights_neg = weights_sort.astype(float)
318     weights_neg[labels_sort == 1] = 0
319
320     total_pos_w = weights_pos.sum()
321     total_neg_w = weights_neg.sum()
322
323     sum_pos = np.cumsum(weights_pos)
324     sum_neg = np.cumsum(weights_neg)
325     error_1 = sum_pos + total_neg_w - sum_neg
326     error_2 = sum_neg + total_pos_w - sum_pos
327
328     min_err_1 = np.min(error_1).astype(float)
329     min_err_2 = np.min(error_2).astype(float)
330     # get the minimum of both, and the minimum between those
331     min_err = np.min([min_err_1, min_err_2])
332
333     if min_err < cls_error:
334         # this is for our best classifier
335         cls_error = min_err
336         if min_err_1 <= min_err_2:
337             polarity = 1
338         else:
339             polarity = 0
340         idx_feature = idxs
341         preds = np.zeros_like(labels_sort)
342         if polarity == 1:
343             threshold = feature_sort[np.argmin(error_1)]
344             preds[feature >= threshold] = 1
345
346         else:
347             threshold = feature_sort[np.argmin(error_2)]
348             preds[feature < threshold] = 1
349
350
351     return [idx_feature, threshold, polarity, cls_error, preds]
352
353 def cascade(x_train_pos, x_train_neg, y_train_pos, y_train_neg, num_iters
354            =5):
355     # this is for the strong classifier
356     cascade_thresholds = []
357     cascade_feature_idxes = []
358     cascade_polarity = []
359     cascade_preds = []
360     cascade_error = []
361     cascade_tfs = []
362     # we initialize the weights here
363     weights_pos = np.ones((y_train_pos.shape[0],1)) * (1/y_train_pos.shape
364     [0])
365     weights_neg = np.ones((y_train_neg.shape[0],1)) * (1/y_train_neg.shape
366     [0])
367     # now we stack everything to use later
368     weights = np.vstack((weights_pos, weights_neg))
369     features = np.vstack((x_train_pos, x_train_neg))
370     labels = np.vstack((y_train_pos[:,None], y_train_neg[:,None]))

```

```

368     final_preds = np.zeros_like(labels)[: ,0]
369     # normalize the weights
370     weights = weights/weights.sum()
371     for idx in range(num_iters):
372         # get each weak classifier and then adjust the weights according
to their trust factors
373         idx_feature, threshold, pol, err, preds = weak_classifier(features,
labels, weights)
374         # use the error from the weak classifier to update the weights
375         eps = err
376         beta = eps / (1 - eps + 1e-16)
377         tf = ((np.log((1 - eps + 1e-16)/(eps + 1e-16))) * 0.5)
378         new_weights = weights * beta**(np.abs(labels[: ,0]-preds))[: ,None]
379         # normalize our new weights
380         new_weights = new_weights/new_weights.sum()
381         # sum to the final predictions, this is just to keep tabs on it
382         final_preds = final_preds + tf*preds
383         cascade_thresholds.append(threshold)
384         cascade_feature_idxes.append(idx_feature)
385         cascade_polarity.append(pol)
386         cascade_preds.append(preds)
387         cascade_error.append(err)
388         cascade_tfs.append(tf)
389
390         # update the weights
391         weights = new_weights
392
393     # get final cascade outputs
394     cascade_tfs_ = np.array(cascade_tfs)
395
396     ths_tf = cascade_tfs_.sum()/2
397
398     final_cascade_preds = np.zeros_like(labels)[: ,0]
399     final_cascade_preds[final_preds >= ths_tf] = 1
400     final_cascade_preds[final_preds < ths_tf] = -1
401
402     # get the negative indexes and positive indexes
403     idx_neg = np.argwhere(labels == 0)[: ,0]
404     idx_pos = np.argwhere(labels == 1)[: ,0]
405     # get the fpr and fnr
406     fpr = (final_cascade_preds[idx_neg] == 1).sum() / y_train_neg.shape[0]
407     fnr = (final_cascade_preds[idx_pos] == 0).sum() / y_train_pos.shape[0]
408     # now we decide which to keep and which to discard,
409     # we only keep the negative samples that have been misclassified
410     neg_preds = final_cascade_preds[idx_neg]
411     idx_keep = np.argwhere(neg_preds == 1)[: ,0]
412     # return these since we will use them for the next strong classifier
413     new_y_train_neg = y_train_neg[idx_keep]
414     new_x_train_neg = x_train_neg[idx_keep]
415
416     strong_classifier = [cascade_thresholds, cascade_feature_idxes,
cascade_polarity, cascade_tfs, cascade_error]
417
418     return new_x_train_neg, new_y_train_neg, strong_classifier, fpr, fnr

```

```
419
420
421 # train cascade
422
423 num_casc = np.arange(1,11,1)
424 fpr_list = []
425 fnr_list = []
426 classifiers = []
427 fpr_cdx = 1
428 fnr_cdx = 1
429 new_x_train_neg = x_train_neg
430 new_y_train_neg = y_train_neg
431 for cdx in num_casc:
432     # we return the new features and subsequent strong classifiers use
    those
433     new_x_train_neg_, new_y_train_neg_, strong_class_stage, fpr, fnr =
    cascade(x_train_pos, new_x_train_neg, y_train_pos, new_y_train_neg,
    num_iters=2)
434     classifiers.append(strong_class_stage)
435     fpr_cdx = fpr_cdx*fpr
436     fnr_cdx = fnr_cdx*fnr
437     fpr_list.append(fpr_cdx)
438     fnr_list.append(fnr_cdx)
439
440     new_x_train_neg = new_x_train_neg_
441     new_y_train_neg = new_y_train_neg_
442     # stop training if we run out of negative samples
443     if len(new_x_train_neg) == 0:
444         break
445
446 classifiers = np.array(classifiers)
447
448 # plot fpr and fnr
449 fig = plt.figure(figsize=(11,11))
450 plt.plot(num_casc, fpr_list, marker="o", label="FPR")
451 plt.plot(num_casc, fnr_list, marker="v", label="FNR")
452
453 plt.xticks(num_casc)
454 plt.legend()
455 plt.savefig(f'imgs/cascade_train.png', bbox_inches='tight', pad_inches=0)
456 plt.close()
457
458 # inference
459 # stack features, labels
460 features_test = np.vstack((x_test_pos, x_test_neg))
461 labels_test = np.vstack((y_test_pos[:,None], y_test_neg[:,None]))
462 fpr_test_list = []
463 fnr_test_list = []
464 fnr_cdx = 1
465 fpr_cdx = 1
466 num_tpos = y_test_pos.shape[0]
467 num_tneg = y_test_neg.shape[0]
468
469 for classifier in classifiers:
```

```
470 # loop through the stages of the cascade
471 preds_test_cascade = np.zeros_like(labels_test)[: ,0]
472 # get the following from each strong classifier
473 feature_idx = classifier[1]
474 thresholds = classifier[0]
475 polarities = classifier[3]
476 trust_factors = classifier[4]
477
478 for cdx in range(feature_idx.shape[0]):
479     # now loop through each weak classifier and start building the
strong classifier output
480     preds_test_classifier = np.zeros_like(labels_test)[: ,0]
481     feature_idx = int(feature_idx[cdx])
482     threshold = thresholds[cdx]
483     polarity = polarities[cdx]
484     trust_factor = trust_factors[cdx]
485
486     feature_test = features_test[:,feature_idx]
487
488     if polarity == 1:
489         preds_test_classifier[feature_test >= threshold] = 1
490         preds_test_classifier[feature_test < threshold] = 0
491     else:
492         preds_test_classifier[feature_test < threshold] = 1
493         preds_test_classifier[feature_test >= threshold] = 0
494
495     preds_test_cascade = preds_test_cascade + preds_test_classifier*
trust_factor
496     # now we adjust the final output
497     ths_tf = trust_factors.sum()/2
498
499     preds_test_cascade[preds_test_cascade >= ths_tf] = 1
500     preds_test_cascade[preds_test_cascade < ths_tf] = 0
501     # get the fpr and fnr
502     idx_pos = np.argwhere(labels_test == 1)[: ,0]
503     idx_neg = np.argwhere(labels_test == 0)[: ,0]
504
505     tot_pos = (labels_test == 1).sum()
506     tot_neg = (labels_test == 0).sum()
507     fpr = (preds_test_cascade[idx_neg] == 1).sum() / num_tneg
508     fnr = (preds_test_cascade[idx_pos] == 0).sum() / num_tpos
509     # keep the misclassified negative samples which means all the positive
predictions
510     idx_keep = np.argwhere(preds_test_cascade == 1)[: ,0]
511
512     features_test = features_test[idx_keep]
513     labels_test = labels_test[idx_keep]
514
515     # cumulative fpr and fnr
516     fnr_cdx = fnr*fnr_cdx
517     fpr_cdx = fpr*fpr_cdx
518     fpr_test_list.append(fpr_cdx)
519     fnr_test_list.append(fnr_cdx)
520
```

```
521 # plot testing fpr and fnr
522 fig = plt.figure(figsize=(11,11))
523 plt.plot(num_casc,fpr_test_list, marker="o",label="FPR")
524 plt.plot(num_casc,fnr_test_list, marker="v",label="FNR")
525
526 plt.xticks(num_casc)
527 plt.legend()
528 plt.savefig(f'imgs/cascade_test.png', bbox_inches='tight', pad_inches=0)
529 plt.close()
```

Listing 1: Source code